


Message Passing and Node Classification

Revisiting Node Classification




- Given labels of some nodes
- Let's predict labels of unlabeled nodes
- This is called semi-supervised node classification

An alternate framework for node classification – Message Passing

- **Intuition: Correlations** exist in networks.
 - In other words: Similar nodes are connected
 - **Key concept** is **collective classification**: Idea of assigning labels to all nodes in a network together
- **We will look at three techniques today:**
 - **Relational classification**
 - **Iterative classification**
 - **Belief propagation**

Correlations in Networks


- Individual behaviors are **correlated** in the network
- **Correlation**: nearby nodes have the same color (belonging to the same class)




Correlations in Networks

- **Main types of dependencies that lead to correlation:**

Homophily




Influence



Homophily

- **Homophily**: The tendency of individuals to associate and bond with similar others
 - *“Birds of a feather flock together”*
 - It has been observed in a vast array of network studies, based on a variety of attributes (e.g., age, gender, organizational role, etc.)
 - **Example**: Researchers who focus on the same research area are more likely to establish a connection (meeting at conferences, interacting in academic talks, etc.)


Homophily



Influence


- **Influence:** Social connections can influence the individual characteristics of a person.
- **Example:** I recommend my musical preferences to my friends, until one of them grows to like my same favorite genres!

Influence



Leveraging the correlation

- How do we leverage this correlation observed in networks to help predict node labels?



How do we predict the labels for the nodes in grey?

Motivation

- **Similar nodes are typically close together or directly connected in the network:**
 - **Guilt-by-association:** If I am connected to a node with label X , then I am likely to have label X as well.
 - **Example: Malicious/benign web page:**
Malicious web pages link to one another to increase visibility, look credible, and rank higher in search engines

Motivation

- **Classification label** of a node v in network may depend on:
 - **Features** of v
 - **Labels** of the nodes in v 's **neighborhood**
 - **Features** of the nodes in v 's **neighborhood**

Collective Classification Overview

- **Intuition:** Simultaneous classification of interlinked nodes using correlations
- Probabilistic framework
- **Markov Assumption:** *the label Y_v of one node v depends on the labels of its neighbors N_v*
$$P(Y_v) = P(Y_v | N_v)$$
- Collective classification involves 3 steps:

Local Classifier

- Assign initial labels

Relational Classifier

- Capture correlations between nodes

Collective Inference

- Propagate correlations through network

Collective Classification Overview

Local Classifier

- Assign initial labels

Local Classifier: Used for initial label assignment

- Predicts label based on node attributes/features
- Standard classification task
- Does not use network information

Relational Classifier

- Capture correlations between nodes

Relational Classifier: Capture correlations

- Learns a classifier to label one node based on the labels and/or attributes of its neighbors
- This is where network information is used

Collective Inference

- Propagate correlations through network


Collective Inference: Propagate the correlation

- Apply relational classifier to each node iteratively
- Iterate until the inconsistency between neighboring labels is minimized
- Network structure affects the final prediction

Problem Setting

- How to predict the labels Y_v for the unlabeled nodes v (in grey color)?
- Each node v has a feature vector f_v
- Labels for some nodes are given (1 for green, 0 for red)
- **Task:** Find $P(Y_v)$ given all features and the network

$$P(Y_v) = ?$$



Intuition and Techniques

- We focus on semi-supervised node classification
- Intuition is based on **homophily**: Similar nodes are typically close together or directly connected
- **Three techniques we will introduce:**
 - **Relational classification**
 - **Iterative classification**
 - **Belief propagation**

Probabilistic Relational Classifier

- **Basic idea:** Class probability Y_v of node v is a weighted average of class probabilities of its neighbors
- For **labeled nodes** v , initialize label Y_v with ground-truth label Y_v^*
- For **unlabeled nodes**, initialize $Y_v = 0.5$
- **Update** all nodes in a random order until convergence or until maximum number of iterations is reached

Probabilistic Relational Classifier

- **Update** for each node v and label c (e.g. 0 or 1)

$$P(Y_v = c) = \frac{1}{\sum_{(v,u) \in E} A_{v,u}} \sum_{(v,u) \in E} A_{v,u} P(Y_u = c)$$


- If edges have strength/weight information, $A_{v,u}$ can be the edge weight between v and u
 - $P(Y_v = c)$ is the probability of node v having label c
- **Challenges:**
 - Convergence is not guaranteed
 - Model cannot use node feature information

Example

Initialization:

- All labeled nodes with their labels
- All unlabeled nodes 0.5 (belonging to class 1 with probability 0.5)

Let $P_{Y_1} = P(Y_1 = 1)$




1st Iteration – Update of node 3

- Update for the 1st Iteration:


- For node 3, $N_3 = \{1, 2, 4\}$

$$P(Y_3) = (0 + 0 + 0.5)/3 = 0.17$$



1st Iteration – Update of Node 4


- Update for the 1st Iteration:
 - For node 4, $N_4 = \{1, 3, 5, 6\}$



After Node 3 is updated


1st Iteration – Update of node 5

- Update for the 1st Iteration:
 - For node 5, $N_5 = \{4, 6, 7, 8\}$




After 1st Iteration

After Iteration 1 (a round of updates for all unlabeled nodes)




After 2nd Iteration

After Iteration 2



Convergence after 4 iterations

- All scores stabilize after 4 iterations. We therefore predict:
 - **Nodes 4, 5, 8, 9 belong to class 1** ($P_{Y_v} > 0.5$)
 - **Nodes 3 belong to class 0** ($P_{Y_v} < 0.5$)



Iterative Classification

- Relational classifiers **do not use node attributes**. How can one leverage them?
- **Main idea of iterative classification:** Classify node v based on its **attributes** f_v as well as **labels** z_v of neighbor set N_v

Iterative Classification Approach

- **Input: Graph**

- f_v : feature vector for node v
- Some nodes v are labeled with Y_v

- **Task:** Predict label of unlabeled nodes


- **Approach: Train two classifiers:**

- $\phi_1(f_v)$ = Predict node label based on node feature vector f_v
- $\phi_2(f_v, z_v)$ = Predict label based on node feature vector f_v and summary z_v of labels of v 's neighbors.

Computing Summary Z_v

How do we compute the summary z_v of labels of v 's neighbors N_v ?

- Ideas: $Z_v = \mathbf{vector}$
 - Histogram of the number (or fraction) of each label in N_v
 - Most common label in N_v
 - Number of different labels in N_v



Iterative Classifier Approach

■ Phase 1: Classify based on node attributes alone

- On a **training set**, train classifier (e.g., linear classifier, neural networks, ...):
- $\phi_1(f_v)$ to predict Y_v based on f_v
- $\phi_2(f_v, z_v)$ to predict Y_v based on f_v and summary z_v of labels of v 's neighbors

■ Phase 2: Iterate till convergence

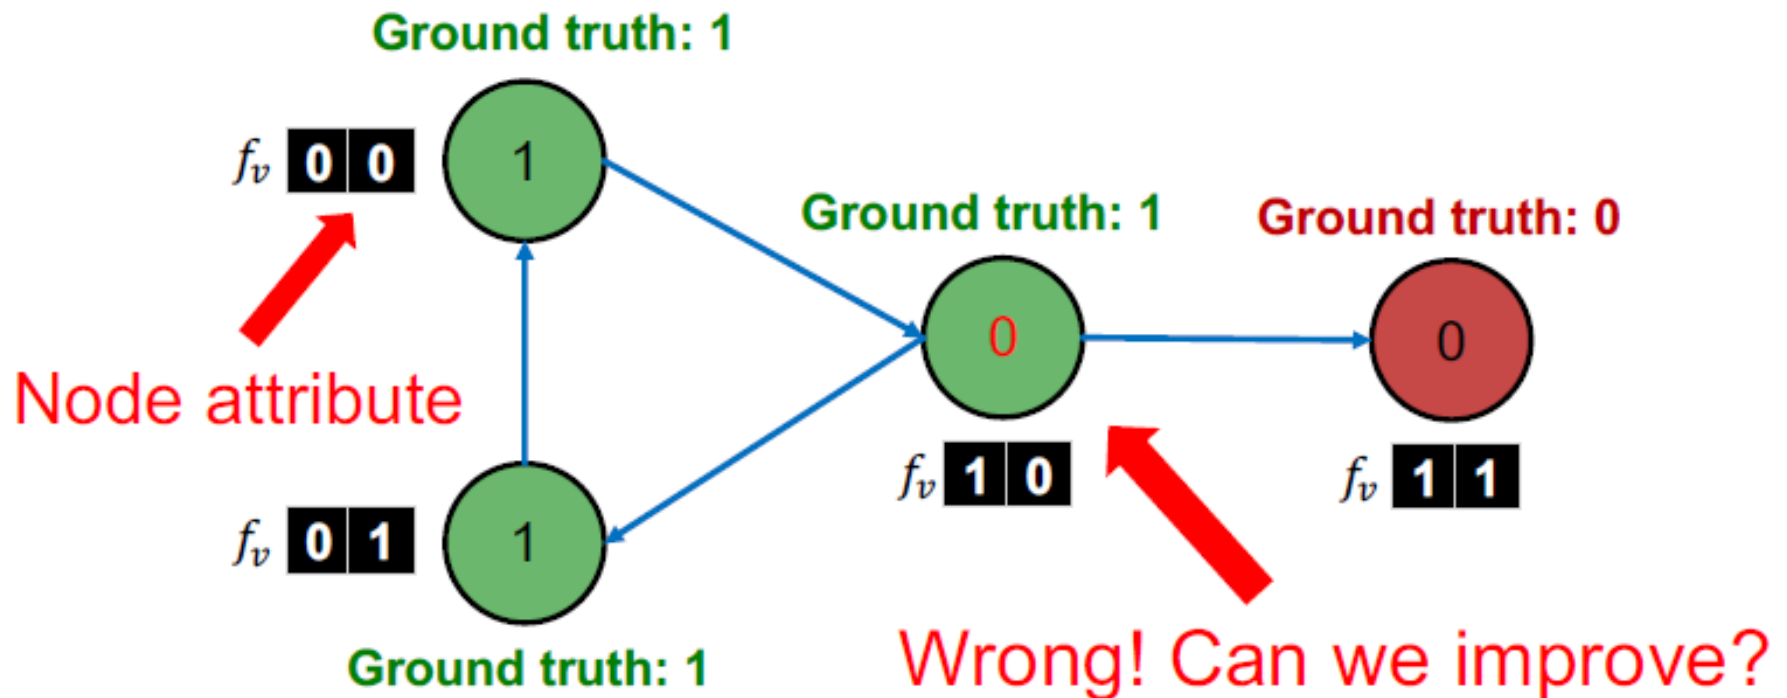
- On **test set**, set labels Y_v based on the classifier ϕ_1 , compute z_v and **predict the labels with ϕ_2**
- **Repeat** for each node v
 - Update z_v based on Y_u for all $u \in N_v$
 - Update Y_v based on the new z_v (ϕ_2)
- Iterate until class labels stabilize or max number of iterations is reached
- Note: Convergence is not guaranteed

Example Web Page Classification

- **Input:** Graph of web pages
- **Node:** Web page
- **Edge:** Hyper-link between web pages
 - **Directed edge:** a page points to another page
- **Node features:** Webpage description
 - For simplicity, we only consider 2 binary features
- **Task:** Predict the topic of the webpage

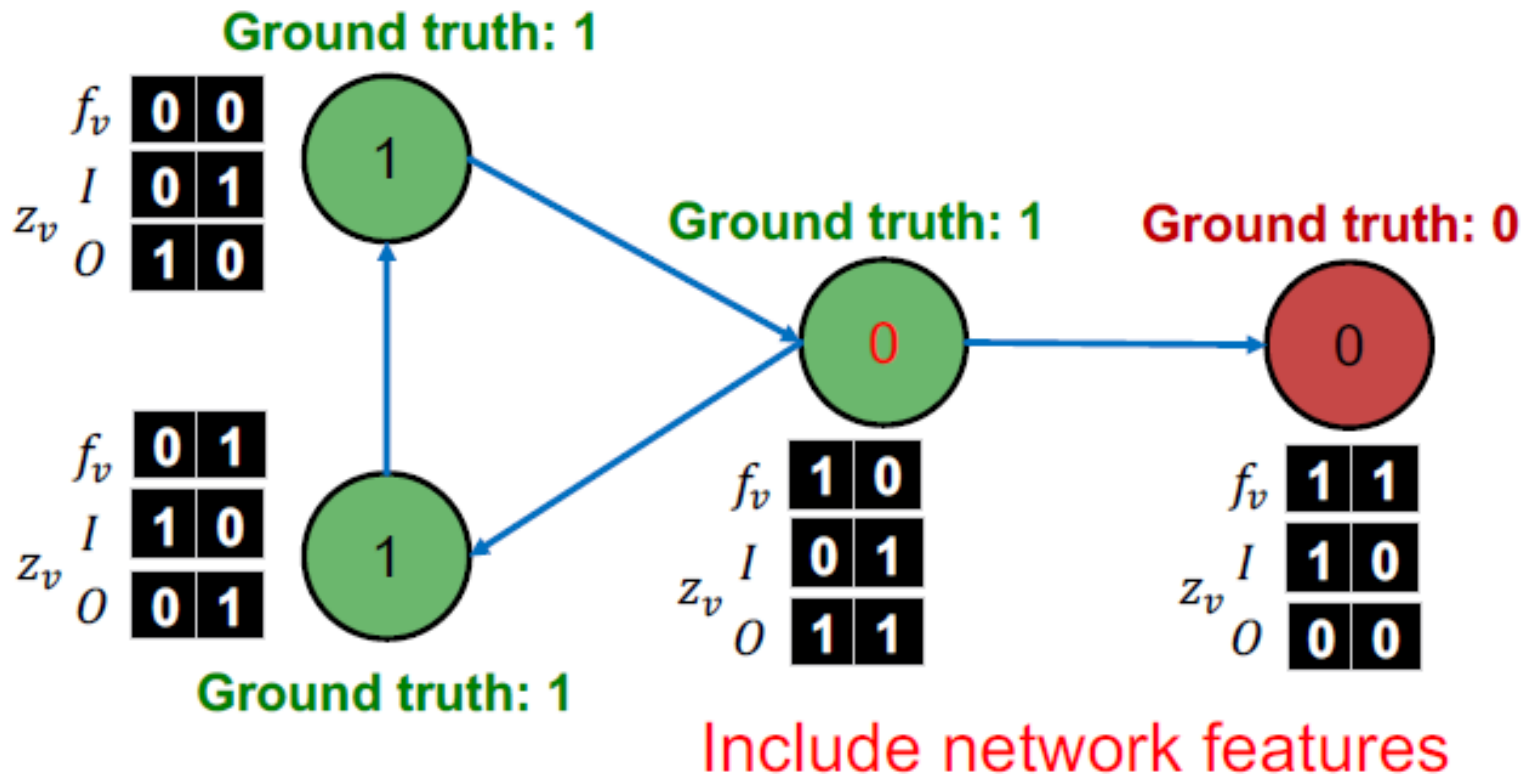
Example Web Page Classification

- **Baseline:** train a classifier (e.g., linear classifier) to classify pages based on binary node attributes.



Example: Web Page Classification


- Each node maintains **vectors \mathbf{z}_v of neighborhood labels**:
 I = **Incoming** neighbor label information vector
 O = **Outgoing** neighbor label information vector
- $I_0 = 1$ if at least one of the incoming pages is labelled 0.
 Similar definitions for I_1 , O_0 , and O_1



Include network features

Iterative Classifier Step 1

- On a different **training set**, train two classifiers:
 - Node attribute vector only (green circles): ϕ_1
 - Node attribute and link vectors (red circles): ϕ_2



- 1.
- 2.
- 3.

Train classifier

Apply classifier to test set

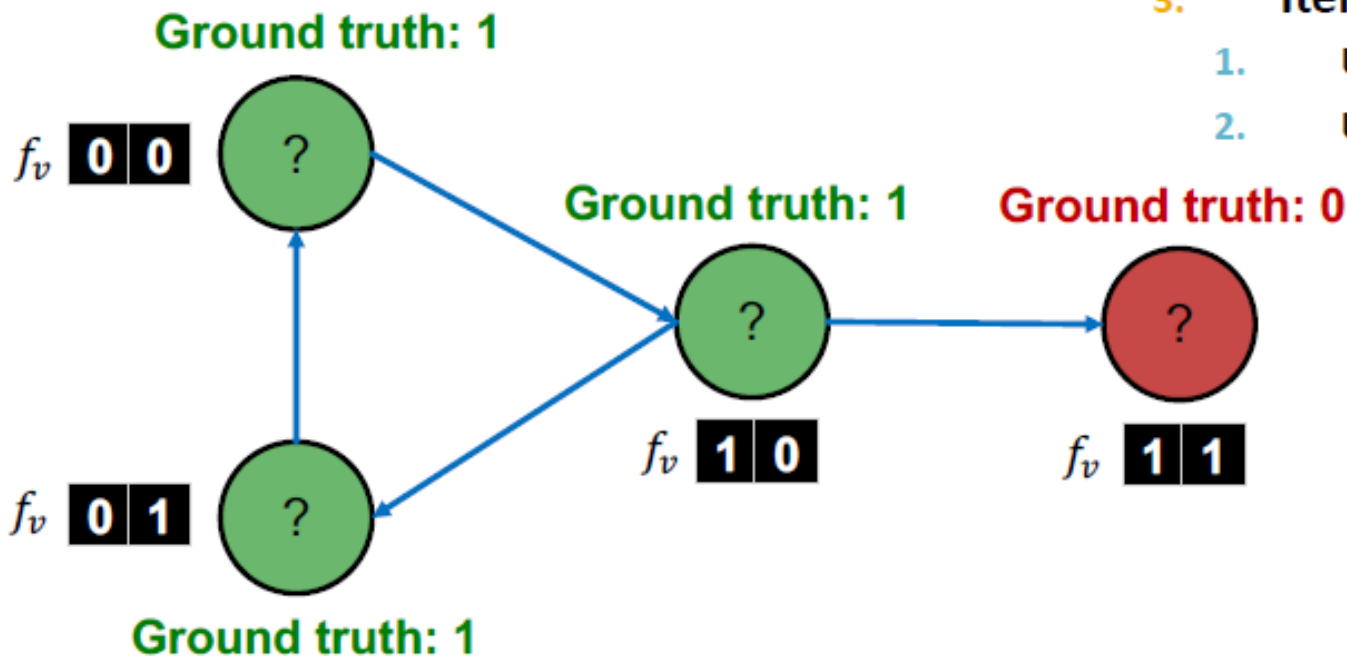
Iterate

1. Update relational features z_v
2. Update label Y_v

Iterative Classifier Step 2

- On the **test set**:

- Use trained node feature vector classifier ϕ_1 to set Y_v




1. Train classifier
2. **Apply classifier to test set**
3. Iterate
 1. Update relational features z_v
 2. Update label Y_v

Iterative Classifier Step 2

- On the **test set**:

- Use trained node feature vector classifier ϕ_1 to set Y_v

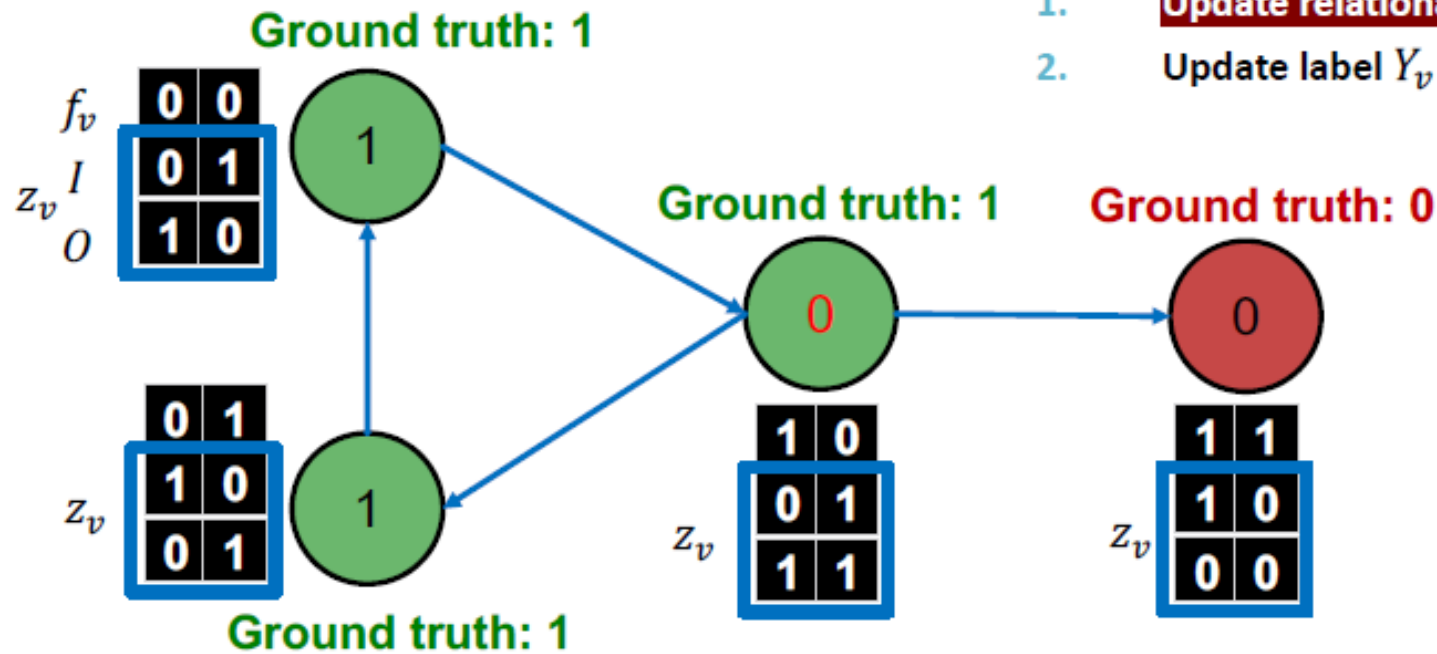
1. Train classifier
2. **Apply classifier to test set**
3. Iterate
 1. Update relational features z_v
 2. Update label Y_v



Iterative Classifier Step 3a

- Update z_v for all nodes


1. Train classifier
2. Apply classifier to test
3. Iterate
 1. **Update relational features z_v**
 2. Update label Y_v



Iterative Classifier Step 3b

- Re-classify all nodes with ϕ_2

1. Train classifier
2. Apply classifier to test
3. Iterate
 1. Update relational features z_v
 2. **Update label** Y_v



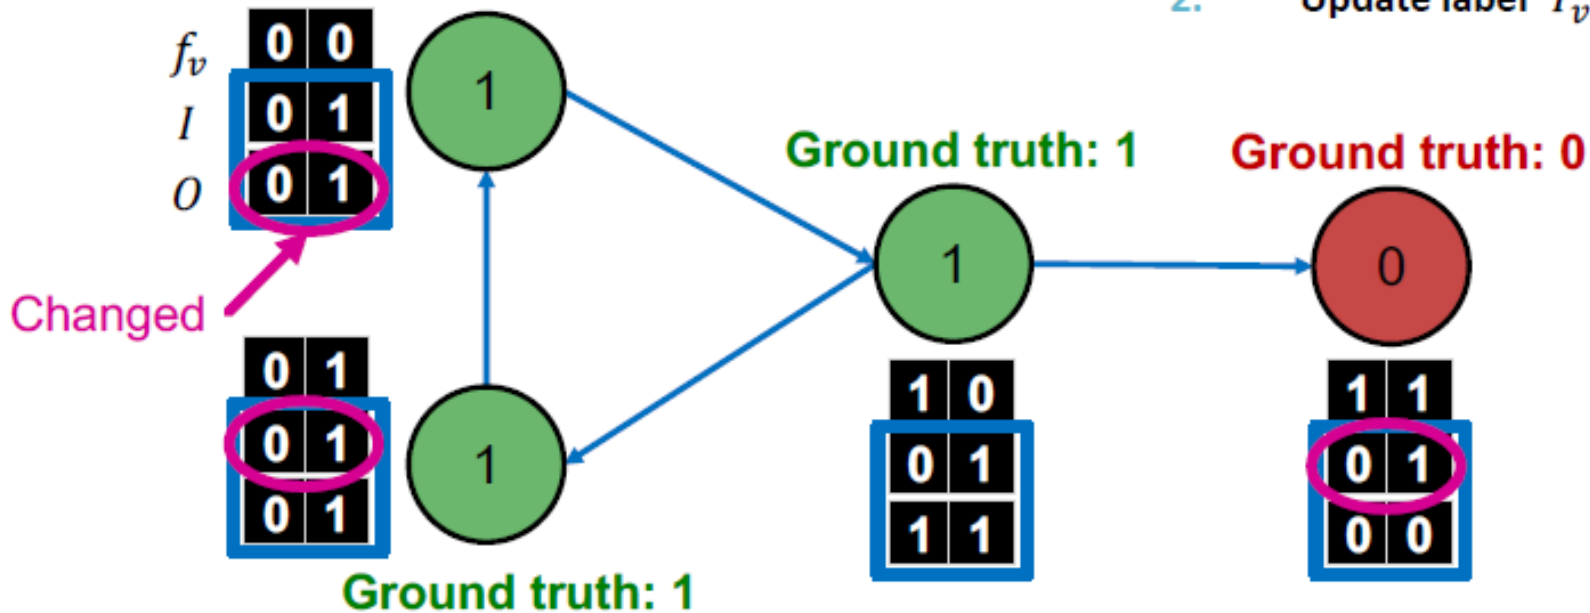
Now it's correct prediction!

Iterative Classifier Step Iterate

- Continue until convergence

- Update z_v
- Update $Y_v = \phi_2(f_v, z_v)$

Ground truth: 1




1. Train classifier
2. Apply classifier to test
3. **Iterate**
 1. Update relational features z_v
 2. Update label Y_v

Iterative Classifier Prediction

- Stop iteration

- After convergence or when maximum iterations are reached

1. Train classifier
2. Apply classifier to test set
3. Iterate
 1. Update relational features z_v
 2. Update label Y_v



Loopy Belief Propagation

- Belief Propagation is a dynamic programming approach to **answering probability queries in a graph** (e.g. probability of node v belonging to class 1)
- Iterative process in which neighbor nodes “talk” to each other, **passing messages**

“I (node v) believe you (node u) belong to class 1 with likelihood ...”



- When **consensus is reached**, calculate final belief

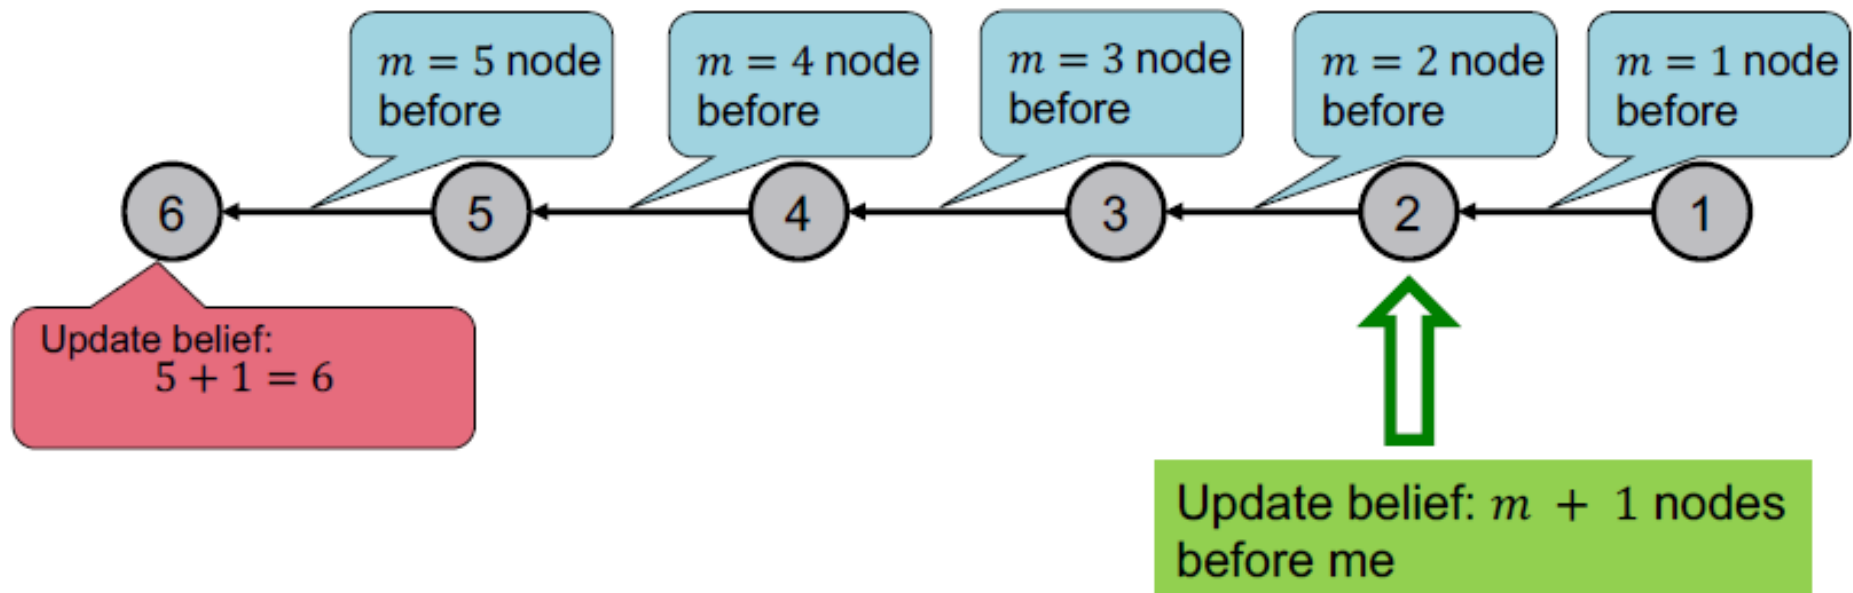
Message Passing Basics

Task: Count the number of nodes in a graph

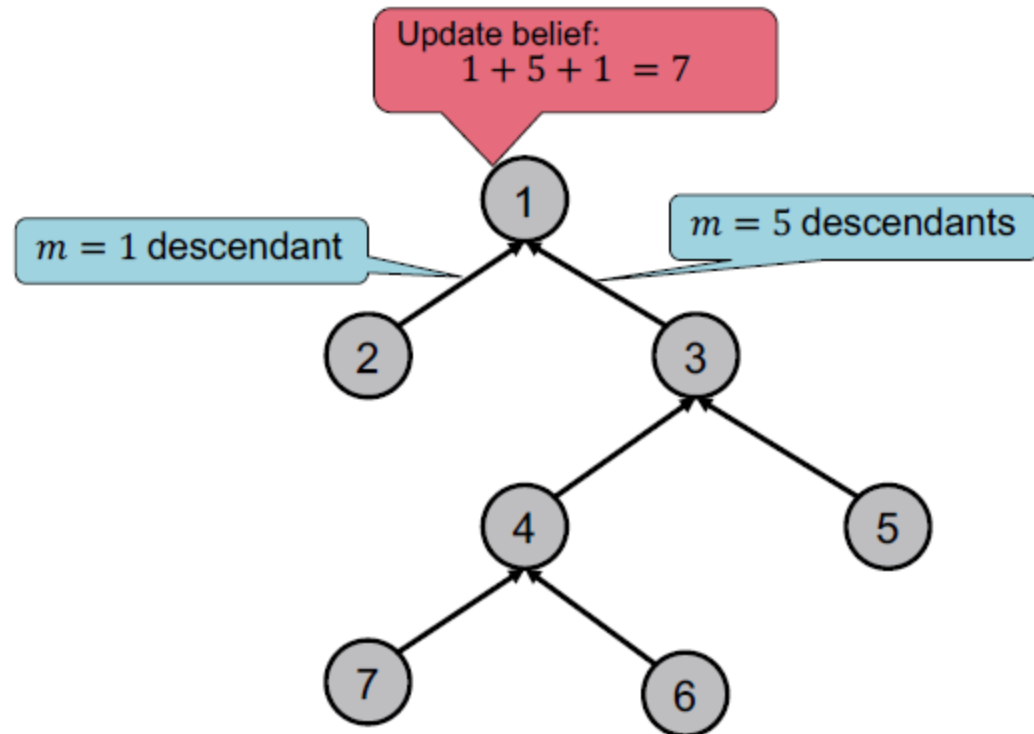

Condition: Each node can only interact (pass message) with its neighbors

Solution: Each node listens to the message from its neighbor, updates it, and passes it forward

m : the message




Belief Propagation in a tree



Loopy Belief Propagation

What message will i send to j ?

- It depends on what i hears from its neighbors
- Each neighbor passes a message to i its beliefs of the state of i




I (node i) believe that
you (node j) belong to
class Y_j with probability
...



Notation

- **Label-label potential matrix ψ** : Dependency between a node and its neighbor. $\psi(Y_i, Y_j)$ is proportional to the probability of a node j being in class Y_j given that it has neighbor i in class Y_i .
- **Prior belief ϕ** : $\phi(Y_i)$ is proportional to the probability of node i being in class Y_i .
- $m_{i \rightarrow j}(Y_j)$ is i 's message / estimate of j being in class Y_j .
- \mathcal{L} is the set of all classes/labels

Loopy Belief Propagation



1. Initialize all messages to 1
2. Repeat for each node:

Label-label
potential

All messages sent by
neighbors from
previous round


$$m_{i \rightarrow j}(Y_j) = \sum_{Y_i \in \mathcal{L}} \psi(Y_i, Y_j) \phi_i(Y_i) \prod_{k \in N_i \setminus j} m_{k \rightarrow i}(Y_i), \forall Y_j \in \mathcal{L}$$

Sum over all states Prior

Loopy Belief Propagation

After convergence:

$b_i(Y_i)$ = node i 's belief of being in class Y_i



All messages from
neighbors

Prior


$$b_i(Y_i) = \phi_i(Y_i) \prod_{j \in N_i} m_{j \rightarrow i}(Y_i), \quad \forall Y_i \in \mathcal{L}$$

Loopy Belief Propagation in Graphs

- Now we consider a graph with cycles
- There is no longer an ordering of nodes
- We apply the same algorithm as in previous slides:
 - Start from arbitrary nodes
 - Follow the edges to update the neighboring nodes

What if the graph has cycles?


What if our graph has cycles?



Messages from different subgraphs are **no longer independent!**

But we can still run BP, but it will pass messages in loops.

What can go wrong?



- **Beliefs may not converge**
 - Message $m_{u \rightarrow i}(Y_i)$ is based on initial belief of i , not a **separate evidence** for i
 - The initial belief of i (which could be incorrect) is reinforced by the cycle $i \rightarrow j \rightarrow k \rightarrow u \rightarrow i$
- However, in practice, Loopy BP is still a good heuristic for complex graphs which contain many branches.