

[spark_starter]

September 12, 2021

1 Working with Apache Spark using pyspark

2 ~ Install spark and pyspark

Install spark on your server

```
sudo apt install spark
```

spark can be accessed from command line using the spark-shell command (here the shell acts as a driver program)

```
spark-shell -c spark.driver.bindAddress=127.0.0.1
```

Install the pyspark package using pip

```
pip install pyspark
```

```
[1]: pip show pyspark
```

```
Name: pyspark
Version: 3.1.2
Summary: Apache Spark Python API
Home-page: https://github.com/apache/spark/tree/master/python
Author: Spark Developers
Author-email: dev@spark.apache.org
License: http://www.apache.org/licenses/LICENSE-2.0
Location: /home/iitp/anaconda3/lib/python3.8/site-packages
Requires: py4j
Required-by:
Note: you may need to restart the kernel to use updated packages.
```

```
[2]: import random, os, shutil
import numpy as np
import matplotlib.pyplot as plt
import pyspark
```

3 Create spark session

a spark session must be created before working with spark.

```
[3]: # we need to start a spark session
from pyspark.sql import SparkSession
from pyspark import SparkContext

spark = SparkSession.builder.appName('Lab6').getOrCreate()
spark
```

```
[3]: <pyspark.sql.session.SparkSession at 0x7fcedbaceaf0>
```

4 The *context* object

```
[4]: sc = spark.sparkContext
print('sparkContext:',sc)
```

```
sparkContext: <SparkContext master=local[*] appName=Lab6>
```

5 RDD - Resilient Distributed Datasets

RDDs are **immutable** collection of datasets that work in parallel read more@
<https://spark.apache.org/docs/latest/rdd-programming-guide.html#resilient-distributed-datasets-rdds>

- **Transformations** are lazy operations on RDDs - stores actions (in a DAG) rather than actual transformation
 - read more@ <https://spark.apache.org/docs/latest/rdd-programming-guide.html#transformations>
- **Actions** on RDDs produce results (using the RDD DAG)
 - read more@ <https://spark.apache.org/docs/latest/rdd-programming-guide.html#actions>

```
[6]: help(pyspark.RDD)
```

Help on class RDD in module pyspark.rdd:

```
class RDD(builtins.object)
|   RDD(jrdd, ctx, jrdd_deserializer=AutoBatchedSerializer(PickleSerializer()))
|
|   A Resilient Distributed Dataset (RDD), the basic abstraction in Spark.
|   Represents an immutable, partitioned collection of elements that can be
|   operated on in parallel.
|
```

```

| Methods defined here:
|
| __add__(self, other)
|     Return the union of this RDD and another one.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize([1, 1, 2, 3])
|     >>> (rdd + rdd).collect()
|     [1, 1, 2, 3, 1, 1, 2, 3]
|
| __getnewargs__(self)
|
| __init__(self, jrdd, ctx,
jrdd_deserializer=AutoBatchedSerializer(PickleSerializer()))
|     Initialize self.  See help(type(self)) for accurate signature.
|
| __repr__(self)
|     Return repr(self).
|
| aggregate(self, zeroValue, seqOp, combOp)
|     Aggregate the elements of each partition, and then the results for all
|     the partitions, using a given combine functions and a neutral "zero
|     value."
|
|     The functions ``op(t1, t2)`` is allowed to modify ``t1`` and return it
|     as its result value to avoid object allocation; however, it should not
|     modify ``t2``.
|
|     The first function (seqOp) can return a different result type, U, than
|     the type of this RDD. Thus, we need one operation for merging a T into
|     an U and one operation for merging two U
|
|     Examples
|     -----
|     >>> seqOp = (lambda x, y: (x[0] + y, x[1] + 1))
|     >>> combOp = (lambda x, y: (x[0] + y[0], x[1] + y[1]))
|     >>> sc.parallelize([1, 2, 3, 4]).aggregate((0, 0), seqOp, combOp)
|     (10, 4)
|     >>> sc.parallelize([]).aggregate((0, 0), seqOp, combOp)
|     (0, 0)
|
| aggregateByKey(self, zeroValue, seqFunc, combFunc, numPartitions=None,
partitionFunc=<function portable_hash at 0x7fcedbed8670>)
|     Aggregate the values of each key, using given combine functions and a
neutral
|     "zero value". This function can return a different result type, U, than
the type

```

| of the values in this RDD, V. Thus, we need one operation for merging a
 V into
 | a U and one operation for merging two U's, The former operation is used
 for merging
 | values within a partition, and the latter is used for merging values
 between
 | partitions. To avoid memory allocation, both of these functions are
 | allowed to modify and return their first argument instead of creating a
 new U.

| barrier(self)
 | Marks the current stage as a barrier stage, where Spark must launch all
 tasks together.

| In case of a task failure, instead of only restarting the failed task,
 Spark will abort the

| entire stage and relaunch all tasks for this stage.

| The barrier execution mode feature is experimental and it only handles
 limited scenarios.

| Please read the linked SPIP and design docs to understand the
 limitations and future plans.

| .. versionadded:: 2.4.0

| Returns

| -----

| :class:`RDDBarrier`

| instance that provides actions within a barrier stage.

| See Also

| -----

| pyspark.BarrierTaskContext

| Notes

| -----

| For additional information see

| - `SPIP: Barrier Execution Mode

<<http://jira.apache.org/jira/browse/SPARK-24374>>`_

| - `Design Doc <<https://jira.apache.org/jira/browse/SPARK-24582>>`_

| This API is experimental

| cache(self)

| Persist this RDD with the default storage level (`MEMORY_ONLY`).

| cartesian(self, other)

| Return the Cartesian product of this RDD and another one, that is, the
 | RDD of all pairs of elements ``(a, b)`` where ``a`` is in ``self`` and

```

|     ``b`` is in `other`.
|
|     Examples
|     -----
|
|     >>> rdd = sc.parallelize([1, 2])
|     >>> sorted(rdd.cartesian(rdd).collect())
|     [(1, 1), (1, 2), (2, 1), (2, 2)]
|
| checkpoint(self)
|     Mark this RDD for checkpointing. It will be saved to a file inside the
|     checkpoint directory set with :meth:`SparkContext.setCheckpointDir` and
|     all references to its parent RDDs will be removed. This function must
|     be called before any job has been executed on this RDD. It is strongly
|     recommended that this RDD is persisted in memory, otherwise saving it
|     on a file will require recomputation.
|
| coalesce(self, numPartitions, shuffle=False)
|     Return a new RDD that is reduced into `numPartitions` partitions.
|
|     Examples
|     -----
|
|     >>> sc.parallelize([1, 2, 3, 4, 5], 3).glom().collect()
|     [[1], [2, 3], [4, 5]]
|     >>> sc.parallelize([1, 2, 3, 4, 5], 3).coalesce(1).glom().collect()
|     [[1, 2, 3, 4, 5]]
|
| cogroup(self, other, numPartitions=None)
|     For each key k in `self` or `other`, return a resulting RDD that
|     contains a tuple with the list of values for that key in `self` as
|     well as `other`.
|
|     Examples
|     -----
|
|     >>> x = sc.parallelize([("a", 1), ("b", 4)])
|     >>> y = sc.parallelize([("a", 2)])
|     >>> [(x, tuple(map(list, y))) for x, y in
sorted(list(x.cogroup(y).collect()))]
|     [('a', ([1], [2])), ('b', ([4], []))]
|
| collect(self)
|     Return a list that contains all of the elements in this RDD.
|
|     Notes
|     ----
|
|     This method should only be used if the resulting array is expected
|     to be small, as all the data is loaded into the driver's memory.
|
| collectAsMap(self)

```

```

|     Return the key-value pairs in this RDD to the master as a dictionary.
|
|     Notes
|     -----
|     This method should only be used if the resulting data is expected
|     to be small, as all the data is loaded into the driver's memory.
|
|     Examples
|     -----
|     >>> m = sc.parallelize([(1, 2), (3, 4)]).collectAsMap()
|     >>> m[1]
|     2
|     >>> m[3]
|     4
|
|     collectWithJobGroup(self, groupId, description, interruptOnCancel=False)
|         When collect rdd, use this method to specify job group.
|
|     .. versionadded:: 3.0.0
|     .. deprecated:: 3.1.0
|         Use :class:`pyspark.InheritableThread` with the pinned thread mode
|     enabled.
|
|     combineByKey(self, createCombiner, mergeValue, mergeCombiners,
| numPartitions=None, partitionFunc=<function portable_hash at 0x7fcedbed8670>)
|         Generic function to combine the elements for each key using a custom
|         set of aggregation functions.
|
|         Turns an RDD[(K, V)] into a result of type RDD[(K, C)], for a "combined
|         type" C.
|
|         Users provide three functions:
|
|             - `createCombiner`, which turns a V into a C (e.g., creates
|               a one-element list)
|             - `mergeValue`, to merge a V into a C (e.g., adds it to the end of
|               a list)
|             - `mergeCombiners`, to combine two C's into a single one (e.g.,
| merges
|               the lists)
|
|         To avoid memory allocation, both mergeValue and mergeCombiners are
| allowed to
|         modify and return their first argument instead of creating a new C.
|
|         In addition, users can control the partitioning of the output RDD.
|
|     Notes

```

V and C can be different -- for example, one might group an RDD of type
(Int, Int) into an RDD of type (Int, List[Int]).

Examples

>>> x = sc.parallelize([("a", 1), ("b", 1), ("a", 2)])
>>> def to_list(a):
... return [a]
...
>>> def append(a, b):
... a.append(b)
... return a
...
>>> def extend(a, b):
... a.extend(b)
... return a
...
>>> sorted(x.combineByKey(to_list, append, extend).collect())
[('a', [1, 2]), ('b', [1])]

count(self)

Return the number of elements in this RDD.

Examples

>>> sc.parallelize([2, 3, 4]).count()
3

countApprox(self, timeout, confidence=0.95)

Approximate version of count() that returns a potentially incomplete
result within a timeout, even if not all tasks have finished.

Examples

>>> rdd = sc.parallelize(range(1000), 10)
>>> rdd.countApprox(1000, 1.0)
1000

countApproxDistinct(self, relativeSD=0.05)

Return approximate number of distinct elements in the RDD.

Parameters

relativeSD : float, optional
Relative accuracy. Smaller values create
counters that require more space.
It must be greater than 0.000017.

```

|
| Notes
| -----
| The algorithm used is based on streamlib's implementation of
| "HyperLogLog in Practice: Algorithmic Engineering of a State
| of The Art Cardinality Estimation Algorithm", available here
| <https://doi.org/10.1145/2452376.2452456>`_.
|
| Examples
| -----
| >>> n = sc.parallelize(range(1000)).map(str).countApproxDistinct()
| >>> 900 < n < 1100
| True
| >>> n = sc.parallelize([i % 20 for i in
range(1000)]).countApproxDistinct()
| >>> 16 < n < 24
| True
|
| countByKey(self)
| Count the number of elements for each key, and return the result to the
| master as a dictionary.
|
| Examples
| -----
| >>> rdd = sc.parallelize([("a", 1), ("b", 1), ("a", 1)])
| >>> sorted(rdd.countByKey().items())
| [('a', 2), ('b', 1)]
|
| countByValue(self)
| Return the count of each unique value in this RDD as a dictionary of
| (value, count) pairs.
|
| Examples
| -----
| >>> sorted(sc.parallelize([1, 2, 1, 2, 2], 2).countByValue().items())
| [(1, 2), (2, 3)]
|
| distinct(self, numPartitions=None)
| Return a new RDD containing the distinct elements in this RDD.
|
| Examples
| -----
| >>> sorted(sc.parallelize([1, 1, 2, 3]).distinct().collect())
| [1, 2, 3]
|
| filter(self, f)
| Return a new RDD containing only the elements that satisfy a predicate.
|

```



```

|     Examples
|     -----
|     >>> rdd = sc.parallelize([1, 2, 3, 4, 5])
|     >>> rdd.filter(lambda x: x % 2 == 0).collect()
|     [2, 4]
|
| first(self)
|     Return the first element in this RDD.
|
|     Examples
|     -----
|     >>> sc.parallelize([2, 3, 4]).first()
|     2
|     >>> sc.parallelize([]).first()
|     Traceback (most recent call last):
|
|         ...
|     ValueError: RDD is empty
|
| flatMap(self, f, preservesPartitioning=False)
|     Return a new RDD by first applying a function to all elements of this
|     RDD, and then flattening the results.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize([2, 3, 4])
|     >>> sorted(rdd.flatMap(lambda x: range(1, x)).collect())
|     [1, 1, 1, 2, 2, 3]
|     >>> sorted(rdd.flatMap(lambda x: [(x, x), (x, x)]).collect())
|     [(2, 2), (2, 2), (3, 3), (3, 3), (4, 4), (4, 4)]
|
| flatMapValues(self, f)
|     Pass each value in the key-value pair RDD through a flatMap function
|     without changing the keys; this also retains the original RDD's
|     partitioning.
|
|     Examples
|     -----
|     >>> x = sc.parallelize([("a", ["x", "y", "z"]), ("b", ["p", "r"])]])
|     >>> def f(x): return x
|     >>> x.flatMapValues(f).collect()
|     [('a', 'x'), ('a', 'y'), ('a', 'z'), ('b', 'p'), ('b', 'r')]
|
| fold(self, zeroValue, op)
|     Aggregate the elements of each partition, and then the results for all
|     the partitions, using a given associative function and a neutral "zero
value."
|
|     The function ``op(t1, t2)`` is allowed to modify ``t1`` and return it

```

as its result value to avoid object allocation; however, it should not modify ``t2``.

This behaves somewhat differently from fold operations implemented for non-distributed collections in functional languages like Scala. This fold operation may be applied to partitions individually, and then fold those results into the final result, rather than apply the fold to each element sequentially in some defined ordering. For functions that are not commutative, the result may differ from that of a fold applied to a non-distributed collection.

Examples

```
>>> from operator import add
>>> sc.parallelize([1, 2, 3, 4, 5]).fold(0, add)
15
```

`foldByKey(self, zeroValue, func, numPartitions=None, partitionFunc=<function portable_hash at 0x7fcedbed8670>)`

Merge the values for each key using an associative function "func" and a neutral "zeroValue" which may be added to the result an arbitrary number of times, and must not change the result (e.g., 0 for addition, or 1 for multiplication.).

Examples

```
>>> rdd = sc.parallelize([("a", 1), ("b", 1), ("a", 1)])
>>> from operator import add
>>> sorted(rdd.foldByKey(0, add).collect())
[('a', 2), ('b', 1)]
```

`foreach(self, f)`

Applies a function to all elements of this RDD.

Examples

```
>>> def f(x): print(x)
>>> sc.parallelize([1, 2, 3, 4, 5]).foreach(f)
```

`foreachPartition(self, f)`

Applies a function to each partition of this RDD.

Examples

```
>>> def f(iterator):
...     for x in iterator:
...         print(x)
>>> sc.parallelize([1, 2, 3, 4, 5]).foreachPartition(f)
```

```

| fullOuterJoin(self, other, numPartitions=None)
|     Perform a right outer join of `self` and `other`.
|
|     For each element (k, v) in `self`, the resulting RDD will either
|     contain all pairs (k, (v, w)) for w in `other`, or the pair
|     (k, (v, None)) if no elements in `other` have key k.
|
|     Similarly, for each element (k, w) in `other`, the resulting RDD will
|     either contain all pairs (k, (v, w)) for v in `self`, or the pair
|     (k, (None, w)) if no elements in `self` have key k.
|
|     Hash-partitions the resulting RDD into the given number of partitions.
|
|     Examples
|     -----
|     >>> x = sc.parallelize([("a", 1), ("b", 4)])
|     >>> y = sc.parallelize([("a", 2), ("c", 8)])
|     >>> sorted(x.fullOuterJoin(y).collect())
|     [('a', (1, 2)), ('b', (4, None)), ('c', (None, 8))]
|
| getCheckpointFile(self)
|     Gets the name of the file to which this RDD was checkpointed
|
|     Not defined if RDD is checkpointed locally.
|
| getNumPartitions(self)
|     Returns the number of partitions in RDD
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize([1, 2, 3, 4], 2)
|     >>> rdd.getNumPartitions()
|     2
|
| getResourceProfile(self)
|     Get the :class:`pyspark.resource.ResourceProfile` specified with this
RDD or None
|     if it wasn't specified.
|
|     .. versionadded:: 3.1.0
|
|     Returns
|     -----
|     :py:class:`pyspark.resource.ResourceProfile`
|         The the user specified profile or None if none were specified
|
|     Notes

```

```

|         -----
|         This API is experimental
|
|     getStorageLevel(self)
|         Get the RDD's current storage level.
|
|         Examples
|         -----
|         >>> rdd1 = sc.parallelize([1,2])
|         >>> rdd1.getStorageLevel()
|         StorageLevel(False, False, False, False, 1)
|         >>> print(rdd1.getStorageLevel())
|         Serialized 1x Replicated
|
|     glom(self)
|         Return an RDD created by coalescing all elements within each partition
|         into a list.
|
|         Examples
|         -----
|         >>> rdd = sc.parallelize([1, 2, 3, 4], 2)
|         >>> sorted(rdd.glom().collect())
|         [[1, 2], [3, 4]]
|
|     groupBy(self, f, numPartitions=None, partitionFunc=<function portable_hash
at 0x7fcedbed8670>)
|         Return an RDD of grouped items.
|
|         Examples
|         -----
|         >>> rdd = sc.parallelize([1, 1, 2, 3, 5, 8])
|         >>> result = rdd.groupBy(lambda x: x % 2).collect()
|         >>> sorted([(x, sorted(y)) for (x, y) in result])
|         [(0, [2, 8]), (1, [1, 1, 3, 5])]
|
|     groupByKey(self, numPartitions=None, partitionFunc=<function portable_hash
at 0x7fcedbed8670>)
|         Group the values for each key in the RDD into a single sequence.
|         Hash-partitions the resulting RDD with numPartitions partitions.
|
|         Notes
|         -----
|         If you are grouping in order to perform an aggregation (such as a
|         sum or average) over each key, using reduceByKey or aggregateByKey will
|         provide much better performance.
|
|         Examples
|         -----

```

```

|     >>> rdd = sc.parallelize([("a", 1), ("b", 1), ("a", 1)])
|     >>> sorted(rdd.groupByKey().mapValues(len).collect())
|     [('a', 2), ('b', 1)]
|     >>> sorted(rdd.groupByKey().mapValues(list).collect())
|     [('a', [1, 1]), ('b', [1])]
|
| groupWith(self, other, *others)
|     Alias for cogroup but with support for multiple RDDs.
|
|     Examples
|     -----
|     >>> w = sc.parallelize([("a", 5), ("b", 6)])
|     >>> x = sc.parallelize([("a", 1), ("b", 4)])
|     >>> y = sc.parallelize([("a", 2)])
|     >>> z = sc.parallelize([("b", 42)])
|     >>> [(x, tuple(map(list, y))) for x, y in sorted(list(w.groupWith(x, y,
z).collect()))]
|     [('a', ([5], [1], [2], [])), ('b', ([6], [4], [], [42]))]
|
| histogram(self, buckets)
|     Compute a histogram using the provided buckets. The buckets
|     are all open to the right except for the last which is closed.
|     e.g. [1,10,20,50] means the buckets are [1,10) [10,20) [20,50],
|     which means  $1 \leq x < 10$ ,  $10 \leq x < 20$ ,  $20 \leq x \leq 50$ . And on the input of 1
|     and 50 we would have a histogram of 1,0,1.
|
|     If your histogram is evenly spaced (e.g. [0, 10, 20, 30]),
|     this can be switched from an  $O(\log n)$  insertion to  $O(1)$  per
|     element (where  $n$  is the number of buckets).
|
|     Buckets must be sorted, not contain any duplicates, and have
|     at least two elements.
|
|     If `buckets` is a number, it will generate buckets which are
|     evenly spaced between the minimum and maximum of the RDD. For
|     example, if the min value is 0 and the max is 100, given `buckets`
|     as 2, the resulting buckets will be [0,50) [50,100]. `buckets` must
|     be at least 1. An exception is raised if the RDD contains infinity.
|     If the elements in the RDD do not vary ( $\max == \min$ ), a single bucket
|     will be used.
|
|     The return value is a tuple of buckets and histogram.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize(range(51))
|     >>> rdd.histogram(2)
|     ([0, 25, 50], [25, 26])

```

```

|     >>> rdd.histogram([0, 5, 25, 50])
|     ([0, 5, 25, 50], [5, 20, 26])
|     >>> rdd.histogram([0, 15, 30, 45, 60]) # evenly spaced buckets
|     ([0, 15, 30, 45, 60], [15, 15, 15, 6])
|     >>> rdd = sc.parallelize(["ab", "ac", "b", "bd", "ef"])
|     >>> rdd.histogram(("a", "b", "c"))
|     (('a', 'b', 'c'), [2, 2])
|
| id(self)
|     A unique ID for this RDD (within its SparkContext).
|
| intersection(self, other)
|     Return the intersection of this RDD and another one. The output will
|     not contain any duplicate elements, even if the input RDDs did.
|
|     Notes
|     ----
|     This method performs a shuffle internally.
|
|     Examples
|     -----
|     >>> rdd1 = sc.parallelize([1, 10, 2, 3, 4, 5])
|     >>> rdd2 = sc.parallelize([1, 6, 2, 3, 7, 8])
|     >>> rdd1.intersection(rdd2).collect()
|     [1, 2, 3]
|
| isCheckpointed(self)
|     Return whether this RDD is checkpointed and materialized, either
reliably or locally.
|
| isEmpty(self)
|     Returns true if and only if the RDD contains no elements at all.
|
|     Notes
|     ----
|     An RDD may be empty even when it has at least 1 partition.
|
|     Examples
|     -----
|     >>> sc.parallelize([]).isEmpty()
|     True
|     >>> sc.parallelize([1]).isEmpty()
|     False
|
| isLocallyCheckpointed(self)
|     Return whether this RDD is marked for local checkpointing.
|
|     Exposed for testing.

```

```

| join(self, other, numPartitions=None)
|     Return an RDD containing all pairs of elements with matching keys in
|     `self` and `other`.
|
|     Each pair of elements will be returned as a (k, (v1, v2)) tuple, where
|     (k, v1) is in `self` and (k, v2) is in `other`.
|
|     Performs a hash join across the cluster.
|
|     Examples
|     -----
|
|     >>> x = sc.parallelize([("a", 1), ("b", 4)])
|     >>> y = sc.parallelize([("a", 2), ("a", 3)])
|     >>> sorted(x.join(y).collect())
|     [('a', (1, 2)), ('a', (1, 3))]
|
| keyBy(self, f)
|     Creates tuples of the elements in this RDD by applying `f`.
|
|     Examples
|     -----
|
|     >>> x = sc.parallelize(range(0,3)).keyBy(lambda x: x*x)
|     >>> y = sc.parallelize(zip(range(0,5), range(0,5)))
|     >>> [(x, list(map(list, y))) for x, y in sorted(x.cogroup(y).collect())]
|     [(0, [[0], [0]]), (1, [[1], [1]]), (2, [[], [2]]), (3, [[], [3]]), (4,
|     [[2], [4]])]
|
| keys(self)
|     Return an RDD with the keys of each tuple.
|
|     Examples
|     -----
|
|     >>> m = sc.parallelize([(1, 2), (3, 4)]).keys()
|     >>> m.collect()
|     [1, 3]
|
| leftOuterJoin(self, other, numPartitions=None)
|     Perform a left outer join of `self` and `other`.
|
|     For each element (k, v) in `self`, the resulting RDD will either
|     contain all pairs (k, (v, w)) for w in `other`, or the pair
|     (k, (v, None)) if no elements in `other` have key k.
|
|     Hash-partitions the resulting RDD into the given number of partitions.
|
|     Examples
|     -----

```

```

|     >>> x = sc.parallelize([("a", 1), ("b", 4)])
|     >>> y = sc.parallelize([("a", 2)])
|     >>> sorted(x.leftOuterJoin(y).collect())
|     [('a', (1, 2)), ('b', (4, None))]
|
|     localCheckpoint(self)
|         Mark this RDD for local checkpointing using Spark's existing caching
layer.
|
|         This method is for users who wish to truncate RDD lineages while
skipping the expensive
|         step of replicating the materialized data in a reliable distributed file
system. This is
|         useful for RDDs with long lineages that need to be truncated
periodically (e.g. GraphX).
|
|         Local checkpointing sacrifices fault-tolerance for performance. In
particular, checkpointed
|         data is written to ephemeral local storage in the executors instead of
to a reliable,
|         fault-tolerant storage. The effect is that if an executor fails during
the computation,
|         the checkpointed data may no longer be accessible, causing an
irrecoverable job failure.
|
|         This is NOT safe to use with dynamic allocation, which removes executors
along
|         with their cached blocks. If you must use both features, you are advised
to set
|         `spark.dynamicAllocation.cachedExecutorIdleTimeout` to a high value.
|
|         The checkpoint directory set through
:meth:`SparkContext.setCheckpointDir` is not used.
|
|     lookup(self, key)
|         Return the list of values in the RDD for key `key`. This operation
|         is done efficiently if the RDD has a known partitioner by only
|         searching the partition that the key maps to.
|
|     Examples
|     -----
|     >>> l = range(1000)
|     >>> rdd = sc.parallelize(zip(l, l), 10)
|     >>> rdd.lookup(42) # slow
|     [42]
|     >>> sorted = rdd.sortByKey()
|     >>> sorted.lookup(42) # fast
|     [42]

```



```

|     >>> sorted.lookup(1024)
|     []
|     >>> rdd2 = sc.parallelize([(('a', 'b'), 'c')]).groupByKey()
|     >>> list(rdd2.lookup(('a', 'b'))[0])
|     ['c']
|
| map(self, f, preservesPartitioning=False)
|     Return a new RDD by applying a function to each element of this RDD.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize(["b", "a", "c"])
|     >>> sorted(rdd.map(lambda x: (x, 1)).collect())
|     [('a', 1), ('b', 1), ('c', 1)]
|
| mapPartitions(self, f, preservesPartitioning=False)
|     Return a new RDD by applying a function to each partition of this RDD.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize([1, 2, 3, 4], 2)
|     >>> def f(iterator): yield sum(iterator)
|     >>> rdd.mapPartitions(f).collect()
|     [3, 7]
|
| mapPartitionsWithIndex(self, f, preservesPartitioning=False)
|     Return a new RDD by applying a function to each partition of this RDD,
|     while tracking the index of the original partition.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize([1, 2, 3, 4], 4)
|     >>> def f(splitIndex, iterator): yield splitIndex
|     >>> rdd.mapPartitionsWithIndex(f).sum()
|     6
|
| mapPartitionsWithSplit(self, f, preservesPartitioning=False)
|     Return a new RDD by applying a function to each partition of this RDD,
|     while tracking the index of the original partition.
|
|     .. deprecated:: 0.9.0
|         use :py:meth:`RDD.mapPartitionsWithIndex` instead.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize([1, 2, 3, 4], 4)
|     >>> def f(splitIndex, iterator): yield splitIndex
|     >>> rdd.mapPartitionsWithSplit(f).sum()

```

```

|         6
|
| mapValues(self, f)
|     Pass each value in the key-value pair RDD through a map function
|     without changing the keys; this also retains the original RDD's
|     partitioning.
|
|     Examples
|     -----
|     >>> x = sc.parallelize([("a", ["apple", "banana", "lemon"]), ("b",
["grapes"])]])
|     >>> def f(x): return len(x)
|     >>> x.mapValues(f).collect()
|     [('a', 3), ('b', 1)]
|
| max(self, key=None)
|     Find the maximum item in this RDD.
|
|     Parameters
|     -----
|     key : function, optional
|         A function used to generate key for comparing
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize([1.0, 5.0, 43.0, 10.0])
|     >>> rdd.max()
|     43.0
|     >>> rdd.max(key=str)
|     5.0
|
| mean(self)
|     Compute the mean of this RDD's elements.
|
|     Examples
|     -----
|     >>> sc.parallelize([1, 2, 3]).mean()
|     2.0
|
| meanApprox(self, timeout, confidence=0.95)
|     Approximate operation to return the mean within a timeout
|     or meet the confidence.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize(range(1000), 10)
|     >>> r = sum(range(1000)) / 1000.0
|     >>> abs(rdd.meanApprox(1000) - r) / r < 0.05

```

```

|     True
|
| min(self, key=None)
|     Find the minimum item in this RDD.
|
|     Parameters
|     -----
|     key : function, optional
|         A function used to generate key for comparing
|
|     Examples
|     -----
|
|     >>> rdd = sc.parallelize([2.0, 5.0, 43.0, 10.0])
|     >>> rdd.min()
|     2.0
|     >>> rdd.min(key=str)
|     10.0
|
| name(self)
|     Return the name of this RDD.
|
| partitionBy(self, numPartitions, partitionFunc=<function portable_hash at
0x7fcedbed8670>)
|     Return a copy of the RDD partitioned using the specified partitioner.
|
|     Examples
|     -----
|
|     >>> pairs = sc.parallelize([1, 2, 3, 4, 2, 4, 1]).map(lambda x: (x, x))
|     >>> sets = pairs.partitionBy(2).glom().collect()
|     >>> len(set(sets[0]).intersection(set(sets[1])))
|     0
|
| persist(self, storageLevel=StorageLevel(False, True, False, False, 1))
|     Set this RDD's storage level to persist its values across operations
|     after the first time it is computed. This can only be used to assign
|     a new storage level if the RDD does not have a storage level set yet.
|     If no storage level is specified defaults to (`MEMORY_ONLY`).
|
|     Examples
|     -----
|
|     >>> rdd = sc.parallelize(["b", "a", "c"])
|     >>> rdd.persist().is_cached
|     True
|
| pipe(self, command, env=None, checkCode=False)
|     Return an RDD created by piping elements to a forked external process.
|
|     Parameters

```

```

|         -----
|         command : str
|             command to run.
|         env : dict, optional
|             environment variables to set.
|         checkCode : bool, optional
|             whether or not to check the return value of the shell command.
|
|         Examples
|         -----
|         >>> sc.parallelize(['1', '2', '', '3']).pipe('cat').collect()
|         ['1', '2', '', '3']
|
| randomSplit(self, weights, seed=None)
|     Randomly splits this RDD with the provided weights.
|
|     weights : list
|         weights for splits, will be normalized if they don't sum to 1
|     seed : int, optional
|         random seed
|
|     Returns
|     -----
|     list
|         split RDDs in a list
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize(range(500), 1)
|     >>> rdd1, rdd2 = rdd.randomSplit([2, 3], 17)
|     >>> len(rdd1.collect() + rdd2.collect())
|     500
|     >>> 150 < rdd1.count() < 250
|     True
|     >>> 250 < rdd2.count() < 350
|     True
|
| reduce(self, f)
|     Reduces the elements of this RDD using the specified commutative and
|     associative binary operator. Currently reduces partitions locally.
|
|     Examples
|     -----
|     >>> from operator import add
|     >>> sc.parallelize([1, 2, 3, 4, 5]).reduce(add)
|     15
|     >>> sc.parallelize((2 for _ in range(10))).map(lambda x:
1).cache().reduce(add)

```

```

|     10
|     >>> sc.parallelize([]).reduce(add)
|     Traceback (most recent call last):
|     ...
|     ValueError: Can not reduce() empty RDD
|
|     reduceByKey(self, func, numPartitions=None, partitionFunc=<function
portable_hash at 0x7fcedbed8670>)
|         Merge the values for each key using an associative and commutative
reduce function.
|
|         This will also perform the merging locally on each mapper before
|         sending results to a reducer, similarly to a "combiner" in MapReduce.
|
|         Output will be partitioned with `numPartitions` partitions, or
|         the default parallelism level if `numPartitions` is not specified.
|         Default partitioner is hash-partition.
|
|         Examples
|         -----
|         >>> from operator import add
|         >>> rdd = sc.parallelize([("a", 1), ("b", 1), ("a", 1)])
|         >>> sorted(rdd.reduceByKey(add).collect())
|         [('a', 2), ('b', 1)]
|
|     reduceByKeyLocally(self, func)
|         Merge the values for each key using an associative and commutative
reduce function, but
|         return the results immediately to the master as a dictionary.
|
|         This will also perform the merging locally on each mapper before
|         sending results to a reducer, similarly to a "combiner" in MapReduce.
|
|         Examples
|         -----
|         >>> from operator import add
|         >>> rdd = sc.parallelize([("a", 1), ("b", 1), ("a", 1)])
|         >>> sorted(rdd.reduceByKeyLocally(add).items())
|         [('a', 2), ('b', 1)]
|
|     repartition(self, numPartitions)
|         Return a new RDD that has exactly numPartitions partitions.
|
|         Can increase or decrease the level of parallelism in this RDD.
|         Internally, this uses a shuffle to redistribute data.
|         If you are decreasing the number of partitions in this RDD, consider
|         using `coalesce`, which can avoid performing a shuffle.

```

```

|     Examples
|     -----
|     >>> rdd = sc.parallelize([1,2,3,4,5,6,7], 4)
|     >>> sorted(rdd.glom().collect())
|     [[1], [2, 3], [4, 5], [6, 7]]
|     >>> len(rdd.repartition(2).glom().collect())
|     2
|     >>> len(rdd.repartition(10).glom().collect())
|     10
|
|     repartitionAndSortWithinPartitions(self, numPartitions=None,
partitionFunc=<function portable_hash at 0x7fcedbed8670>, ascending=True,
keyfunc=<function RDD.<lambda> at 0x7fcedbcfc160>)
|         Repartition the RDD according to the given partitioner and, within each
resulting partition,
|         sort records by their keys.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize([(0, 5), (3, 8), (2, 6), (0, 8), (3, 8), (1,
3)])
|     >>> rdd2 = rdd.repartitionAndSortWithinPartitions(2, lambda x: x % 2,
True)
|     >>> rdd2.glom().collect()
|     [[(0, 5), (0, 8), (2, 6)], [(1, 3), (3, 8), (3, 8)]]
|
|     rightOuterJoin(self, other, numPartitions=None)
|         Perform a right outer join of `self` and `other`.
|
|         For each element (k, w) in `other`, the resulting RDD will either
|         contain all pairs (k, (v, w)) for v in this, or the pair (k, (None, w))
|         if no elements in `self` have key k.
|
|         Hash-partitions the resulting RDD into the given number of partitions.
|
|     Examples
|     -----
|     >>> x = sc.parallelize([("a", 1), ("b", 4)])
|     >>> y = sc.parallelize([("a", 2)])
|     >>> sorted(y.rightOuterJoin(x).collect())
|     [('a', (2, 1)), ('b', (None, 4))]
|
|     sample(self, withReplacement, fraction, seed=None)
|         Return a sampled subset of this RDD.
|
|     Parameters
|     -----
|     withReplacement : bool

```

```

|         can elements be sampled multiple times (replaced when sampled out)
|     fraction : float
|         expected size of the sample as a fraction of this RDD's size
|         without replacement: probability that each element is chosen;
fraction must be [0, 1]
|         with replacement: expected number of times each element is chosen;
fraction must be >= 0
|     seed : int, optional
|         seed for the random number generator
|
|     Notes
|     ----
|     This is not guaranteed to provide exactly the fraction specified of the
total
|     count of the given :class:`DataFrame`.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize(range(100), 4)
|     >>> 6 <= rdd.sample(False, 0.1, 81).count() <= 14
|     True
|
|     sampleByKey(self, withReplacement, fractions, seed=None)
|     Return a subset of this RDD sampled by key (via stratified sampling).
|     Create a sample of this RDD using variable sampling rates for
|     different keys as specified by fractions, a key to sampling rate map.
|
|     Examples
|     -----
|     >>> fractions = {"a": 0.2, "b": 0.1}
|     >>> rdd =
sc.parallelize(fractions.keys()).cartesian(sc.parallelize(range(0, 1000)))
|     >>> sample = dict(rdd.sampleByKey(False, fractions,
2).groupByKey().collect())
|     >>> 100 < len(sample["a"]) < 300 and 50 < len(sample["b"]) < 150
|     True
|     >>> max(sample["a"]) <= 999 and min(sample["a"]) >= 0
|     True
|     >>> max(sample["b"]) <= 999 and min(sample["b"]) >= 0
|     True
|
|     sampleStdev(self)
|     Compute the sample standard deviation of this RDD's elements (which
|     corrects for bias in estimating the standard deviation by dividing by
|     N-1 instead of N).
|
|     Examples
|     -----

```

```

|     >>> sc.parallelize([1, 2, 3]).sampleStdev()
|     1.0
|
|     sampleVariance(self)
|         Compute the sample variance of this RDD's elements (which corrects
|         for bias in estimating the variance by dividing by N-1 instead of N).
|
|         Examples
|         -----
|         >>> sc.parallelize([1, 2, 3]).sampleVariance()
|         1.0
|
|     saveAsHadoopDataset(self, conf, keyConverter=None, valueConverter=None)
|         Output a Python RDD of key-value pairs (of form ``RDD[(K, V)]``) to any
Hadoop file
|         system, using the old Hadoop OutputFormat API (mapred package).
Keys/values are
|         converted for output using either user specified converters or, by
default,
|         "org.apache.spark.api.python.JavaToWritableConverter".
|
|         Parameters
|         -----
|         conf : dict
|             Hadoop job configuration
|         keyConverter : str, optional
|             fully qualified classname of key converter (None by default)
|         valueConverter : str, optional
|             fully qualified classname of value converter (None by default)
|
|     saveAsHadoopFile(self, path, outputFormatClass, keyClass=None,
valueClass=None, keyConverter=None, valueConverter=None, conf=None,
compressionCodecClass=None)
|         Output a Python RDD of key-value pairs (of form ``RDD[(K, V)]``) to any
Hadoop file
|         system, using the old Hadoop OutputFormat API (mapred package). Key and
value types
|         will be inferred if not specified. Keys and values are converted for
output using either
|         user specified converters or
"org.apache.spark.api.python.JavaToWritableConverter". The
|         `conf` is applied on top of the base Hadoop conf associated with the
SparkContext
|         of this RDD to create a merged Hadoop MapReduce job configuration for
saving the data.
|
|         Parameters
|         -----

```



```

|     path : str
|         path to Hadoop file
|     outputFormatClass : str
|         fully qualified classname of Hadoop OutputFormat
|         (e.g. "org.apache.hadoop.mapred.SequenceFileOutputFormat")
|     keyClass : str, optional
|         fully qualified classname of key Writable class
|         (e.g. "org.apache.hadoop.io.IntWritable", None by default)
|     valueClass : str, optional
|         fully qualified classname of value Writable class
|         (e.g. "org.apache.hadoop.io.Text", None by default)
|     keyConverter : str, optional
|         fully qualified classname of key converter (None by default)
|     valueConverter : str, optional
|         fully qualified classname of value converter (None by default)
|     conf : dict, optional
|         (None by default)
|     compressionCodecClass : str
|         fully qualified classname of the compression codec class
|         i.e. "org.apache.hadoop.io.compress.GzipCodec" (None by default)
|
|     saveAsNewAPIHadoopDataset(self, conf, keyConverter=None,
valueConverter=None)
|         Output a Python RDD of key-value pairs (of form ``RDD[(K, V)]``) to any
Hadoop file
|         system, using the new Hadoop OutputFormat API (mapreduce package).
Keys/values are
|         converted for output using either user specified converters or, by
default,
|         "org.apache.spark.api.python.JavaToWritableConverter".
|
|     Parameters
|     -----
|     conf : dict
|         Hadoop job configuration
|     keyConverter : str, optional
|         fully qualified classname of key converter (None by default)
|     valueConverter : str, optional
|         fully qualified classname of value converter (None by default)
|
|     saveAsNewAPIHadoopFile(self, path, outputFormatClass, keyClass=None,
valueClass=None, keyConverter=None, valueConverter=None, conf=None)
|         Output a Python RDD of key-value pairs (of form ``RDD[(K, V)]``) to any
Hadoop file
|         system, using the new Hadoop OutputFormat API (mapreduce package). Key
and value types
|         will be inferred if not specified. Keys and values are converted for
output using either

```

```

|         user specified converters or
"org.apache.spark.api.python.JavaToWritableConverter". The
|         `conf` is applied on top of the base Hadoop conf associated with the
SparkContext
|         of this RDD to create a merged Hadoop MapReduce job configuration for
saving the data.
|
|         path : str
|             path to Hadoop file
|         outputFormatClass : str
|             fully qualified classname of Hadoop OutputFormat
|             (e.g.
"org.apache.hadoop.mapreduce.lib.output.SequenceFileOutputFormat")
|         keyClass : str, optional
|             fully qualified classname of key Writable class
|             (e.g. "org.apache.hadoop.io.IntWritable", None by default)
|         valueClass : str, optional
|             fully qualified classname of value Writable class
|             (e.g. "org.apache.hadoop.io.Text", None by default)
|         keyConverter : str, optional
|             fully qualified classname of key converter (None by default)
|         valueConverter : str, optional
|             fully qualified classname of value converter (None by default)
|         conf : dict, optional
|             Hadoop job configuration (None by default)
|
|         saveAsPickleFile(self, path, batchSize=10)
|             Save this RDD as a SequenceFile of serialized objects. The serializer
|             used is :class:`pyspark.serializers.PickleSerializer`, default batch
size
|             is 10.
|
|         Examples
|         -----
|
|         >>> from tempfile import NamedTemporaryFile
|         >>> tmpFile = NamedTemporaryFile(delete=True)
|         >>> tmpFile.close()
|         >>> sc.parallelize([1, 2, 'spark',
'rdd']).saveAsPickleFile(tmpFile.name, 3)
|         >>> sorted(sc.pickleFile(tmpFile.name, 5).map(str).collect())
|         ['1', '2', 'rdd', 'spark']
|
|         saveAsSequenceFile(self, path, compressionCodecClass=None)
|             Output a Python RDD of key-value pairs (of form ``RDD[(K, V)]``) to any
Hadoop file
|             system, using the "org.apache.hadoop.io.Writable" types that we convert
from the
|             RDD's key and value types. The mechanism is as follows:

```

|
| 1. Pyrolite is used to convert pickled Python RDD into RDD of Java
objects.

| 2. Keys and values of this Java RDD are converted to Writables and
written out.

| Parameters

| -----

| path : str

| path to sequence file

| compressionCodecClass : str, optional

| fully qualified classname of the compression codec class

| i.e. "org.apache.hadoop.io.compress.GzipCodec" (None by default)

| saveAsTextFile(self, path, compressionCodecClass=None)

| Save this RDD as a text file, using string representations of elements.

| Parameters

| -----

| path : str

| path to text file

| compressionCodecClass : str, optional

| fully qualified classname of the compression codec class

| i.e. "org.apache.hadoop.io.compress.GzipCodec" (None by default)

| Examples

| -----

| >>> from tempfile import NamedTemporaryFile

| >>> tempFile = NamedTemporaryFile(delete=True)

| >>> tempFile.close()

| >>> sc.parallelize(range(10)).saveAsTextFile(tempFile.name)

| >>> from fileinput import input

| >>> from glob import glob

| >>> ''.join(sorted(input(glob(tempFile.name + "/part-0000*"))))

| '0\n1\n2\n3\n4\n5\n6\n7\n8\n9\n'

| Empty lines are tolerated when saving to text files.

| >>> from tempfile import NamedTemporaryFile

| >>> tempFile2 = NamedTemporaryFile(delete=True)

| >>> tempFile2.close()

| >>> sc.parallelize(['', 'foo', '', 'bar',
'']).saveAsTextFile(tempFile2.name)

| >>> ''.join(sorted(input(glob(tempFile2.name + "/part-0000*"))))

| '\n\nnbar\nfoo\n'

| Using compressionCodecClass

```

|     >>> from tempfile import NamedTemporaryFile
|     >>> tempFile3 = NamedTemporaryFile(delete=True)
|     >>> tempFile3.close()
|     >>> codec = "org.apache.hadoop.io.compress.GzipCodec"
|     >>> sc.parallelize(['foo', 'bar']).saveAsTextFile(tempFile3.name, codec)
|     >>> from fileinput import input, hook_compressed
|     >>> result = sorted(input(glob(tempFile3.name + "/part*.gz"),
openhook=hook_compressed))
|     >>> b''.join(result).decode('utf-8')
|     'bar\nfoo\n'
|
|     setName(self, name)
|         Assign a name to this RDD.
|
|     Examples
|     -----
|     >>> rdd1 = sc.parallelize([1, 2])
|     >>> rdd1.setName('RDD1').name()
|     'RDD1'
|
|     sortBy(self, keyfunc, ascending=True, numPartitions=None)
|         Sorts this RDD by the given keyfunc
|
|     Examples
|     -----
|     >>> tmp = [('a', 1), ('b', 2), ('1', 3), ('d', 4), ('2', 5)]
|     >>> sc.parallelize(tmp).sortBy(lambda x: x[0]).collect()
|     [('1', 3), ('2', 5), ('a', 1), ('b', 2), ('d', 4)]
|     >>> sc.parallelize(tmp).sortBy(lambda x: x[1]).collect()
|     [('a', 1), ('b', 2), ('1', 3), ('d', 4), ('2', 5)]
|
|     sortByKey(self, ascending=True, numPartitions=None, keyfunc=<function
RDD.<lambda> at 0x7fcedbcfc280>)
|         Sorts this RDD, which is assumed to consist of (key, value) pairs.
|
|     Examples
|     -----
|     >>> tmp = [('a', 1), ('b', 2), ('1', 3), ('d', 4), ('2', 5)]
|     >>> sc.parallelize(tmp).sortByKey().first()
|     ('1', 3)
|     >>> sc.parallelize(tmp).sortByKey(True, 1).collect()
|     [('1', 3), ('2', 5), ('a', 1), ('b', 2), ('d', 4)]
|     >>> sc.parallelize(tmp).sortByKey(True, 2).collect()
|     [('1', 3), ('2', 5), ('a', 1), ('b', 2), ('d', 4)]
|     >>> tmp2 = [('Mary', 1), ('had', 2), ('a', 3), ('little', 4), ('lamb',
5)]
|     >>> tmp2.extend([('whose', 6), ('fleece', 7), ('was', 8), ('white', 9)])
|     >>> sc.parallelize(tmp2).sortByKey(True, 3, keyfunc=lambda k:

```

```

k.lower()).collect()
|      [('a', 3), ('fleece', 7), ('had', 2), ('lamb', 5),...('white', 9),
('whose', 6)]
|
| stats(self)
|      Return a :class:`StatCounter` object that captures the mean, variance
|      and count of the RDD's elements in one operation.
|
| stdev(self)
|      Compute the standard deviation of this RDD's elements.
|
|      Examples
|      -----
|      >>> sc.parallelize([1, 2, 3]).stdev()
|      0.816...
|
| subtract(self, other, numPartitions=None)
|      Return each value in `self` that is not contained in `other`.
|
|      Examples
|      -----
|      >>> x = sc.parallelize([("a", 1), ("b", 4), ("b", 5), ("a", 3)])
|      >>> y = sc.parallelize([("a", 3), ("c", None)])
|      >>> sorted(x.subtract(y).collect())
|      [('a', 1), ('b', 4), ('b', 5)]
|
| subtractByKey(self, other, numPartitions=None)
|      Return each (key, value) pair in `self` that has no pair with matching
|      key in `other`.
|
|      Examples
|      -----
|      >>> x = sc.parallelize([("a", 1), ("b", 4), ("b", 5), ("a", 2)])
|      >>> y = sc.parallelize([("a", 3), ("c", None)])
|      >>> sorted(x.subtractByKey(y).collect())
|      [('b', 4), ('b', 5)]
|
| sum(self)
|      Add up the elements in this RDD.
|
|      Examples
|      -----
|      >>> sc.parallelize([1.0, 2.0, 3.0]).sum()
|      6.0
|
| sumApprox(self, timeout, confidence=0.95)
|      Approximate operation to return the sum within a timeout
|      or meet the confidence.

```

```

|
| Examples
| -----
| >>> rdd = sc.parallelize(range(1000), 10)
| >>> r = sum(range(1000))
| >>> abs(rdd.sumApprox(1000) - r) / r < 0.05
| True
|
| take(self, num)
|     Take the first num elements of the RDD.
|
| It works by first scanning one partition, and use the results from
| that partition to estimate the number of additional partitions needed
| to satisfy the limit.
|
| Translated from the Scala implementation in RDD#take().
|
| Notes
| -----
| This method should only be used if the resulting array is expected
| to be small, as all the data is loaded into the driver's memory.
|
| Examples
| -----
| >>> sc.parallelize([2, 3, 4, 5, 6]).cache().take(2)
| [2, 3]
| >>> sc.parallelize([2, 3, 4, 5, 6]).take(10)
| [2, 3, 4, 5, 6]
| >>> sc.parallelize(range(100), 100).filter(lambda x: x > 90).take(3)
| [91, 92, 93]
|
| takeOrdered(self, num, key=None)
|     Get the N elements from an RDD ordered in ascending order or as
|     specified by the optional key function.
|
| Notes
| -----
| This method should only be used if the resulting array is expected
| to be small, as all the data is loaded into the driver's memory.
|
| Examples
| -----
| >>> sc.parallelize([10, 1, 2, 9, 3, 4, 5, 6, 7]).takeOrdered(6)
| [1, 2, 3, 4, 5, 6]
| >>> sc.parallelize([10, 1, 2, 9, 3, 4, 5, 6, 7], 2).takeOrdered(6,
key=lambda x: -x)
| [10, 9, 7, 6, 5, 4]
|

```

```

| takeSample(self, withReplacement, num, seed=None)
|     Return a fixed-size sampled subset of this RDD.
|
|     Notes
|     -----
|     This method should only be used if the resulting array is expected
|     to be small, as all the data is loaded into the driver's memory.
|
|     Examples
|     -----
|     >>> rdd = sc.parallelize(range(0, 10))
|     >>> len(rdd.takeSample(True, 20, 1))
|     20
|     >>> len(rdd.takeSample(False, 5, 2))
|     5
|     >>> len(rdd.takeSample(False, 15, 3))
|     10
|
| toDF(self, schema=None, sampleRatio=None)
|     Converts current :class:`RDD` into a :class:`DataFrame`
|
|     This is a shorthand for ``spark.createDataFrame(rdd, schema,
sampleRatio)``
|
|     Parameters
|     -----
|     schema : :class:`pyspark.sql.types.DataType`, str or list, optional
|         a :class:`pyspark.sql.types.DataType` or a datatype string or a list
of
|         column names, default is None. The data type string format equals
to
|         :class:`pyspark.sql.types.DataType.simpleString`, except that top
level struct type can
|         omit the ``struct<>`` and atomic types use ``typeName()`` as their
format, e.g. use
|         ``byte`` instead of ``tinyint`` for
:class:`pyspark.sql.types.ByteType`.
|         We can also use ``int`` as a short name for
:class:`pyspark.sql.types.IntegerType`.
|         sampleRatio : float, optional
|             the sample ratio of rows used for inferring
|
|     Returns
|     -----
|     :class:`DataFrame`
|
|     Examples
|     -----

```

```

|     >>> rdd.toDF().collect()
|     [Row(name='Alice', age=1)]
|
| toDebugString(self)
|     A description of this RDD and its recursive dependencies for debugging.
|
| toLocalIterator(self, prefetchPartitions=False)
|     Return an iterator that contains all of the elements in this RDD.
|     The iterator will consume as much memory as the largest partition in
this RDD.
|     With prefetch it may consume up to the memory of the 2 largest
partitions.
|
| Parameters
| -----
| prefetchPartitions : bool, optional
|     If Spark should pre-fetch the next partition
|     before it is needed.
|
| Examples
| -----
| >>> rdd = sc.parallelize(range(10))
| >>> [x for x in rdd.toLocalIterator()]
| [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
|
| top(self, num, key=None)
|     Get the top N elements from an RDD.
|
| Notes
| -----
| This method should only be used if the resulting array is expected
| to be small, as all the data is loaded into the driver's memory.
|
| It returns the list sorted in descending order.
|
| Examples
| -----
| >>> sc.parallelize([10, 4, 2, 12, 3]).top(1)
| [12]
| >>> sc.parallelize([2, 3, 4, 5, 6], 2).top(2)
| [6, 5]
| >>> sc.parallelize([10, 4, 2, 12, 3]).top(3, key=str)
| [4, 3, 2]
|
| treeAggregate(self, zeroValue, seqOp, combOp, depth=2)
|     Aggregates the elements of this RDD in a multi-level tree
|     pattern.

```



```
depth : int, optional
    suggested depth of the tree (default: 2)
```

Examples

```
-----
>>> add = lambda x, y: x + y
>>> rdd = sc.parallelize([-5, -4, -3, -2, -1, 1, 2, 3, 4], 10)
>>> rdd.treeAggregate(0, add, add)
-5
>>> rdd.treeAggregate(0, add, add, 1)
-5
>>> rdd.treeAggregate(0, add, add, 2)
-5
>>> rdd.treeAggregate(0, add, add, 5)
-5
>>> rdd.treeAggregate(0, add, add, 10)
-5
```

```
treeReduce(self, f, depth=2)
    Reduces the elements of this RDD in a multi-level tree pattern.
```

Parameters

```
-----
f : function
depth : int, optional
    suggested depth of the tree (default: 2)
```

Examples

```
-----
>>> add = lambda x, y: x + y
>>> rdd = sc.parallelize([-5, -4, -3, -2, -1, 1, 2, 3, 4], 10)
>>> rdd.treeReduce(add)
-5
>>> rdd.treeReduce(add, 1)
-5
>>> rdd.treeReduce(add, 2)
-5
>>> rdd.treeReduce(add, 5)
-5
>>> rdd.treeReduce(add, 10)
-5
```

```
union(self, other)
    Return the union of this RDD and another one.
```

Examples

```
-----
>>> rdd = sc.parallelize([1, 1, 2, 3])
```

```

|     >>> rdd.union(rdd).collect()
|     [1, 1, 2, 3, 1, 1, 2, 3]
|
| unpersist(self, blocking=False)
|     Mark the RDD as non-persistent, and remove all blocks for it from
|     memory and disk.
|
|     .. versionchanged:: 3.0.0
|         Added optional argument `blocking` to specify whether to block until
all     blocks are deleted.
|
| values(self)
|     Return an RDD with the values of each tuple.
|
|     Examples
|     -----
|     >>> m = sc.parallelize([(1, 2), (3, 4)]).values()
|     >>> m.collect()
|     [2, 4]
|
| variance(self)
|     Compute the variance of this RDD's elements.
|
|     Examples
|     -----
|     >>> sc.parallelize([1, 2, 3]).variance()
|     0.666...
|
| withResources(self, profile)
|     Specify a :class:`pyspark.resource.ResourceProfile` to use when
calculating this RDD.
|     This is only supported on certain cluster managers and currently
requires dynamic
|     allocation to be enabled. It will result in new executors with the
resources specified
|     being acquired to calculate the RDD.
|
|     .. versionadded:: 3.1.0
|
|     Notes
|     ----
|     This API is experimental
|
| zip(self, other)
|     Zips this RDD with another one, returning key-value pairs with the
|     first element in each RDD second element in each RDD, etc. Assumes
|     that the two RDDs have the same number of partitions and the same

```

```

|     number of elements in each partition (e.g. one was made through
|     a map on the other).
|
|     Examples
|     -----
|     >>> x = sc.parallelize(range(0,5))
|     >>> y = sc.parallelize(range(1000, 1005))
|     >>> x.zip(y).collect()
|     [(0, 1000), (1, 1001), (2, 1002), (3, 1003), (4, 1004)]
|
| zipWithIndex(self)
|     Zips this RDD with its element indices.
|
|     The ordering is first based on the partition index and then the
|     ordering of items within each partition. So the first item in
|     the first partition gets index 0, and the last item in the last
|     partition receives the largest index.
|
|     This method needs to trigger a spark job when this RDD contains
|     more than one partitions.
|
|     Examples
|     -----
|     >>> sc.parallelize(["a", "b", "c", "d"], 3).zipWithIndex().collect()
|     [('a', 0), ('b', 1), ('c', 2), ('d', 3)]
|
| zipWithUniqueId(self)
|     Zips this RDD with generated unique Long ids.
|
|     Items in the kth partition will get ids k, n+k, 2*n+k, ..., where
|     n is the number of partitions. So there may exist gaps, but this
|     method won't trigger a spark job, which is different from
|     :meth:`zipWithIndex`.
|
|     Examples
|     -----
|     >>> sc.parallelize(["a", "b", "c", "d", "e"],
3).zipWithUniqueId().collect()
|     [('a', 0), ('b', 1), ('c', 4), ('d', 2), ('e', 5)]
|
| -----
| Readonly properties defined here:
|
| context
|     The :class:`SparkContext` that this RDD was created on.
|
| -----
| Data descriptors defined here:

```

```
|
|  __dict__
|      dictionary for instance variables (if defined)
|
|  __weakref__
|      list of weak references to the object (if defined)
```

6 Creating RDDs

6.1 creating from sc.parallelize()

```
[7]: RDD_array1 = sc.parallelize(np.random.randint(0,10,size=10))
```

6.2 apply transformation and actions

```
[8]: RDD_array2 = RDD_array1.map(lambda x: x*2) # 'map' is a transformation
```

```
[9]: a1 = RDD_array1.collect() # 'collect' is an action
a2 = RDD_array2.collect()
print( a1,  a2)
```

```
[7, 9, 2, 0, 4, 1, 2, 0, 5, 0] [14, 18, 4, 0, 8, 2, 4, 0, 10, 0]
```

```
[10]: a3 = RDD_array2.reduce(lambda x,y: x+y) # 'reduce' is an action
print(a3)
```

```
60
```

6.3 creating from files or external objects

```
[11]: RDD_F = sc.textFile("Apple_stock.csv")
```

```
[12]: cf = RDD_F.count() # 'count' is an action
print('Count:',cf)

data = RDD_F.collect() # 'collect' is an action
print(type(data), len(data))

samples = RDD_F.takeSample(True, 3) # takeSample is an action
print('Samples\n', type(samples), len(samples), samples)

first = RDD_F.first()
```

```
print('Header:',first) #<---- shall be used for schema
```

Count: 1597

<class 'list'> 1597

Samples

```
<class 'list'> 3 ['2015-11-16,28.559999465942383,27.75,27.844999313354492,28.54
5000076293945,152426800.0,26.322452545166016', '2011-01-03,11.795000076293945,11
.601428985595703,11.630000114440918,11.770357131958008,445138400.0,10.1062202453
61328', '2012-09-10,24.403213500976562,23.64642906188965,24.301786422729492,23.6
69286727905273,487998000.0,20.410099029541016']
```

Header: Date,High,Low,Open,Close,Volume,Adj Close

7 mapper function used to map values to data frame

```
[13]: def mapper(xs):
        global first
        if xs!=first:
            x = xs.split(",")
            return [str(x[0]), float(x[1]), float(x[2]), float(x[3]), float(x[4]),
↪float(x[5]), float(x[6])]
        else:
            pass
```

7.1 create data frame using mapper

```
[14]: dataset = RDD_F.map(mapper).filter(lambda x: x!=None).toDF(first.split(','))
print(type(dataset), dataset.count())
dataset.printSchema()
dataset.show() # or use .describe()
```

<class 'pyspark.sql.dataframe.DataFrame'> 1596

root

```
|-- Date: string (nullable = true)
|-- High: double (nullable = true)
|-- Low: double (nullable = true)
|-- Open: double (nullable = true)
|-- Close: double (nullable = true)
|-- Volume: double (nullable = true)
|-- Adj Close: double (nullable = true)
```

```
+-----+-----+-----+-----+-----+
--+-----+-----+
|      Date|      High|      Low|      Open|
Close|      Volume|      Adj Close|
```

```

+-----+-----+-----+-----+-----+
--+-----+-----+
|2010-08-02|9.378213882446289|9.272143363952637|
9.30142879486084|9.351785659790039|4.280556E8| 8.029596328735352|
|2010-08-03|9.402142524719238|9.265000343322754|9.321785926818848|9.354642868041
992|4.176536E8| 8.032052993774414|
|2010-08-04|9.438570976257324|9.296786308288574|
9.3871431350708|9.392143249511719|4.203752E8| 8.064249038696289|
|2010-08-05|9.399286270141602|9.305356979370117|
9.34749984741211|9.346428871154785|2.890972E8| 8.024996757507324|
|2010-08-06|9.338929176330566|9.201070785522461|9.277856826782227|9.288928985595
703|4.448976E8| 7.975627899169922|
|2010-08-09|9.362500190734863|9.270357131958008|9.338570594787598|9.348214149475
098| 3.03128E8| 8.026529312133789|
|2010-08-10|9.301786422729492|9.198213577270508|9.280357360839844|9.264642715454
102| 4.5192E8| 7.954774856567383|
|2010-08-11|9.131786346435547|8.921786308288574|9.121429443359375|8.935357093811
035|6.200544E8| 7.672046184539795|
|2010-08-12|9.039285659790039|8.789999961853027|8.810357093811035|8.992500305175
781|5.349204E8|7.7211103439331055|
|2010-08-13| 8.99571418762207|8.896071434020996|8.987500190734863|8.896429061889
648|3.548692E8| 7.638622283935547|
|2010-08-16|8.928929328918457|8.807856559753418|8.842143058776855|
8.84428596496582| 3.1843E8|7.5938496589660645|
|2010-08-17|9.093929290771484|8.899999618530273|8.931428909301758|8.998929023742
676|4.226404E8| 7.726629257202148|
|2010-08-18|9.095356941223145|8.984999656677246|9.012857437133789|9.038213729858
398| 3.39696E8|7.7603607177734375|
|2010-08-19|9.052857398986816|8.881428718566895|9.029999732971191|8.924285888671
875| 4.26706E8| 7.662538051605225|
|2010-08-20|9.068571090698242| 8.89285659790039|
8.90678596496582|8.915714263916016| 3.8423E8| 7.655179500579834|
|2010-08-23| 9.0|8.758929252624512|8.992500305175781|8.778571128845
215|4.140416E8|7.5374250411987305|
|2010-08-24|8.678570747375488|8.523214340209961|8.666786193847656|8.568928718566
895|6.025656E8| 7.357422828674316|
|2010-08-25|8.713929176330566|8.471428871154785|8.501428604125977|8.674642562866
211|5.968676E8| 7.448192596435547|
|2010-08-26|8.776785850524902|8.581428527832031|8.766071319580078|8.581428527832
031|4.665052E8| 7.368154048919678|
|2010-08-27|8.664643287658691|8.412857055664062|8.633929252624512|
8.62928581237793|5.483912E8| 7.409246921539307|
+-----+-----+-----+-----+-----+
--+-----+-----+
only showing top 20 rows

```

8 Linear Regression Example


Not using spark.mllib

```
[15]: # a line is  $y = mx + c$  where  $m$  and  $c$  are the params
ground_m, ground_c = random.randint(0,5), random.randint(0,5) # the ground truth

def line(x):
    global ground_m, ground_c
    return ground_m*x + ground_c

# generate at least 100 samples and save to a file
rdd_x = sc.parallelize(np.random.randint(0,1000,size=100))
rdd_y = rdd_x.map(line)

# generate input data and save to file
data_x, data_y = rdd_x.collect(), rdd_y.collect() #<-- 'collect' is an action
plt.scatter(data_x, data_y)
plt.show()
```



8.1 Estimate using keras

```
[16]: from tensorflow.keras import Model # always use this approach instead of
      ↪ Sequential
      from tensorflow.keras.layers import Input, Dense
      from tensorflow.keras.optimizers import Adam, SGD, RMSprop

      # create a single layer model
      input_layer = Input((1,))
      output_layer = Dense(1)(input_layer)
      model = Model(inputs=input_layer, outputs=output_layer)
      model.compile(loss='mse', optimizer=Adam(learning_rate=0.1), metrics=[])
      model.summary()
```

Model: "model"


Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 1)]	0
dense (Dense)	(None, 1)	2

Total params: 2
Trainable params: 2
Non-trainable params: 0

9 Train model

```
[17]: lossv=[]
      for ep in range(100):
          #print(ep+1)
          hist=model.fit(x=np.array(rdd_x.collect()), #<--- new data every time
          ↪collect() is called
                          y=np.array(rdd_y.collect()),
                          batch_size=32, epochs=20,verbose=0)
          lossv.extend(hist.history['loss'])
      plt.plot(lossv)
```

```
[17]: [<matplotlib.lines.Line2D at 0x7fce94424c40>]
```

10 Check for convergence

check if estimated parameters are close enough to ground truth

```
[18]: model_w8 = model.get_weights()
      est_m, est_c = model_w8[0][0], model_w8[1][0]
      print('Model Estimate:', est_m, est_c)
      print('Ground Truth:', '['+str(ground_m)+']', ground_c)

      print('Delta:', ground_m-est_m, ground_c-est_c)
```

```
Model Estimate: [3.] 3.9999743
Ground Truth: [3] 4
Delta: [0.] 2.574920654296875e-05
```

11 Using spark DataFrames in general

12 Reading Data

```
[19]: dataset = spark.read.option('header',
                                'true' #<---- the csv has a header
                                ).csv("Apple_stock.csv",
                                       inferSchema=True) # add infer schema to load in
                                ↪proper data type

print(type(dataset), dataset.count())
dataset.printSchema()
dataset.show() # or use .describe()
```

```
<class 'pyspark.sql.dataframe.DataFrame'> 1596
```

```
root
```

```
|-- Date: string (nullable = true)
|-- High: double (nullable = true)
|-- Low: double (nullable = true)
|-- Open: double (nullable = true)
|-- Close: double (nullable = true)
|-- Volume: double (nullable = true)
|-- Adj Close: double (nullable = true)
```

```
+-----+-----+-----+-----+
--++-----+
|      Date|      High|      Low|      Open|
Close|      Volume|      Adj Close|
+-----+-----+-----+-----+
--++-----+
|2010-08-02|9.378213882446289|9.272143363952637|
9.30142879486084|9.351785659790039|4.280556E8| 8.029596328735352|
|2010-08-03|9.402142524719238|9.265000343322754|9.321785926818848|9.354642868041
992|4.176536E8| 8.032052993774414|
|2010-08-04|9.438570976257324|9.296786308288574|
9.3871431350708|9.392143249511719|4.203752E8| 8.064249038696289|
|2010-08-05|9.399286270141602|9.305356979370117|
9.34749984741211|9.346428871154785|2.890972E8| 8.024996757507324|
|2010-08-06|9.338929176330566|9.201070785522461|9.277856826782227|9.288928985595
703|4.448976E8| 7.975627899169922|
|2010-08-09|9.362500190734863|9.270357131958008|9.338570594787598|9.348214149475
098| 3.03128E8| 8.026529312133789|
|2010-08-10|9.301786422729492|9.198213577270508|9.280357360839844|9.264642715454
102| 4.5192E8| 7.954774856567383|
|2010-08-11|9.131786346435547|8.921786308288574|9.121429443359375|8.935357093811
035|6.200544E8| 7.672046184539795|
```

```

|2010-08-12|9.039285659790039|8.789999961853027|8.810357093811035|8.992500305175
781|5.349204E8|7.7211103439331055|
|2010-08-13| 8.99571418762207|8.896071434020996|8.987500190734863|8.896429061889
648|3.548692E8| 7.638622283935547|
|2010-08-16|8.928929328918457|8.807856559753418|8.842143058776855|
8.84428596496582| 3.1843E8|7.5938496589660645|
|2010-08-17|9.093929290771484|8.899999618530273|8.931428909301758|8.998929023742
676|4.226404E8| 7.726629257202148|
|2010-08-18|9.095356941223145|8.984999656677246|9.012857437133789|9.038213729858
398| 3.39696E8|7.7603607177734375|
|2010-08-19|9.052857398986816|8.881428718566895|9.029999732971191|8.924285888671
875| 4.26706E8| 7.662538051605225|
|2010-08-20|9.068571090698242| 8.89285659790039|
8.90678596496582|8.915714263916016| 3.8423E8| 7.655179500579834|
|2010-08-23| 9.0|8.758929252624512|8.992500305175781|8.778571128845
215|4.140416E8|7.5374250411987305|
|2010-08-24|8.678570747375488|8.523214340209961|8.666786193847656|8.568928718566
895|6.025656E8| 7.357422828674316|
|2010-08-25|8.713929176330566|8.471428871154785|8.501428604125977|8.674642562866
211|5.968676E8| 7.448192596435547|
|2010-08-26|8.776785850524902|8.581428527832031|8.766071319580078|8.581428527832
031|4.665052E8| 7.368154048919678|
|2010-08-27|8.664643287658691|8.412857055664062|8.633929252624512|
8.62928581237793|5.483912E8| 7.409246921539307|
+-----+-----+-----+-----+-----+
+-----+
only showing top 20 rows

```

```
[20]: dataset.select(['Date', 'Volume']).show()
```

```

+-----+-----+
|      Date|      Volume|
+-----+-----+
|2010-08-02|4.280556E8|
|2010-08-03|4.176536E8|
|2010-08-04|4.203752E8|
|2010-08-05|2.890972E8|
|2010-08-06|4.448976E8|
|2010-08-09| 3.03128E8|
|2010-08-10| 4.5192E8|
|2010-08-11|6.200544E8|
|2010-08-12|5.349204E8|
|2010-08-13|3.548692E8|
|2010-08-16| 3.1843E8|
|2010-08-17|4.226404E8|
|2010-08-18| 3.39696E8|
|2010-08-19| 4.26706E8|

```

```
|2010-08-20| 3.8423E8|
|2010-08-23|4.140416E8|
|2010-08-24|6.025656E8|
|2010-08-25|5.968676E8|
|2010-08-26|4.665052E8|
|2010-08-27|5.483912E8|
+-----+-----+
only showing top 20 rows
```

13 Filter Operation

```
[21]: dataset.filter("Volume>=400000000").select(["Date","Volume","Adj Close"]).show()
```

```
+-----+-----+-----+
|      Date|      Volume|      Adj Close|
+-----+-----+-----+
|2010-08-02|4.280556E8| 8.029596328735352|
|2010-08-03|4.176536E8| 8.032052993774414|
|2010-08-04|4.203752E8| 8.064249038696289|
|2010-08-06|4.448976E8| 7.975627899169922|
|2010-08-10| 4.5192E8| 7.954774856567383|
|2010-08-11|6.200544E8| 7.672046184539795|
|2010-08-12|5.349204E8|7.7211103439331055|
|2010-08-17|4.226404E8| 7.726629257202148|
|2010-08-19| 4.26706E8| 7.662538051605225|
|2010-08-23|4.140416E8|7.5374250411987305|
|2010-08-24|6.025656E8| 7.357422828674316|
|2010-08-25|5.968676E8| 7.448192596435547|
|2010-08-26|4.665052E8| 7.368154048919678|
|2010-08-27|5.483912E8| 7.409246921539307|
|2010-08-31|4.207868E8|7.4546308517456055|
|2010-09-01|6.970376E8|7.6763386726379395|
|2010-09-02|4.154276E8| 7.732760429382324|
|2010-09-03|5.207888E8| 7.935146808624268|
|2010-09-08|5.265512E8| 8.062408447265625|
|2010-09-09|4.385752E8| 8.067007064819336|
+-----+-----+-----+
only showing top 20 rows
```