

# CQL

## (Cassandra Query Language)



**Dr. Rajiv Misra**

**Dept. of Computer Science & Engg.**  
**Indian Institute of Technology Patna**  
**[rajivm@iitp.ac.in](mailto:rajivm@iitp.ac.in)**

# Preface

## Content of this Lecture:

- In this lecture, we will discuss CQL (Cassandra Query Language) Mapping to Cassandra's Internal Data Structure.

# What Problems does CQL Solve?

- **The Awesomeness that is Cassandra:**
  - Distributed columnar data store
  - No single point of failure
  - Optimized for availability (through “Tunably” consistent
  - Optimized for writes
  - Easily maintainable
  - Almost infinitely scalable

# What Problems does CQL Solve? (Contd.)

- **Cassandra's usability challenges**

- NoSQL: “Where are my JOINS? No Schema? De-normalize!?”
- BigTable: “Tables with millions of columns!?”


- **CQL Saves the day!**

- A *best-practices* interface to Cassandra
- Uses familiar SQL-Like language


# C\* Data Model

Keyspace

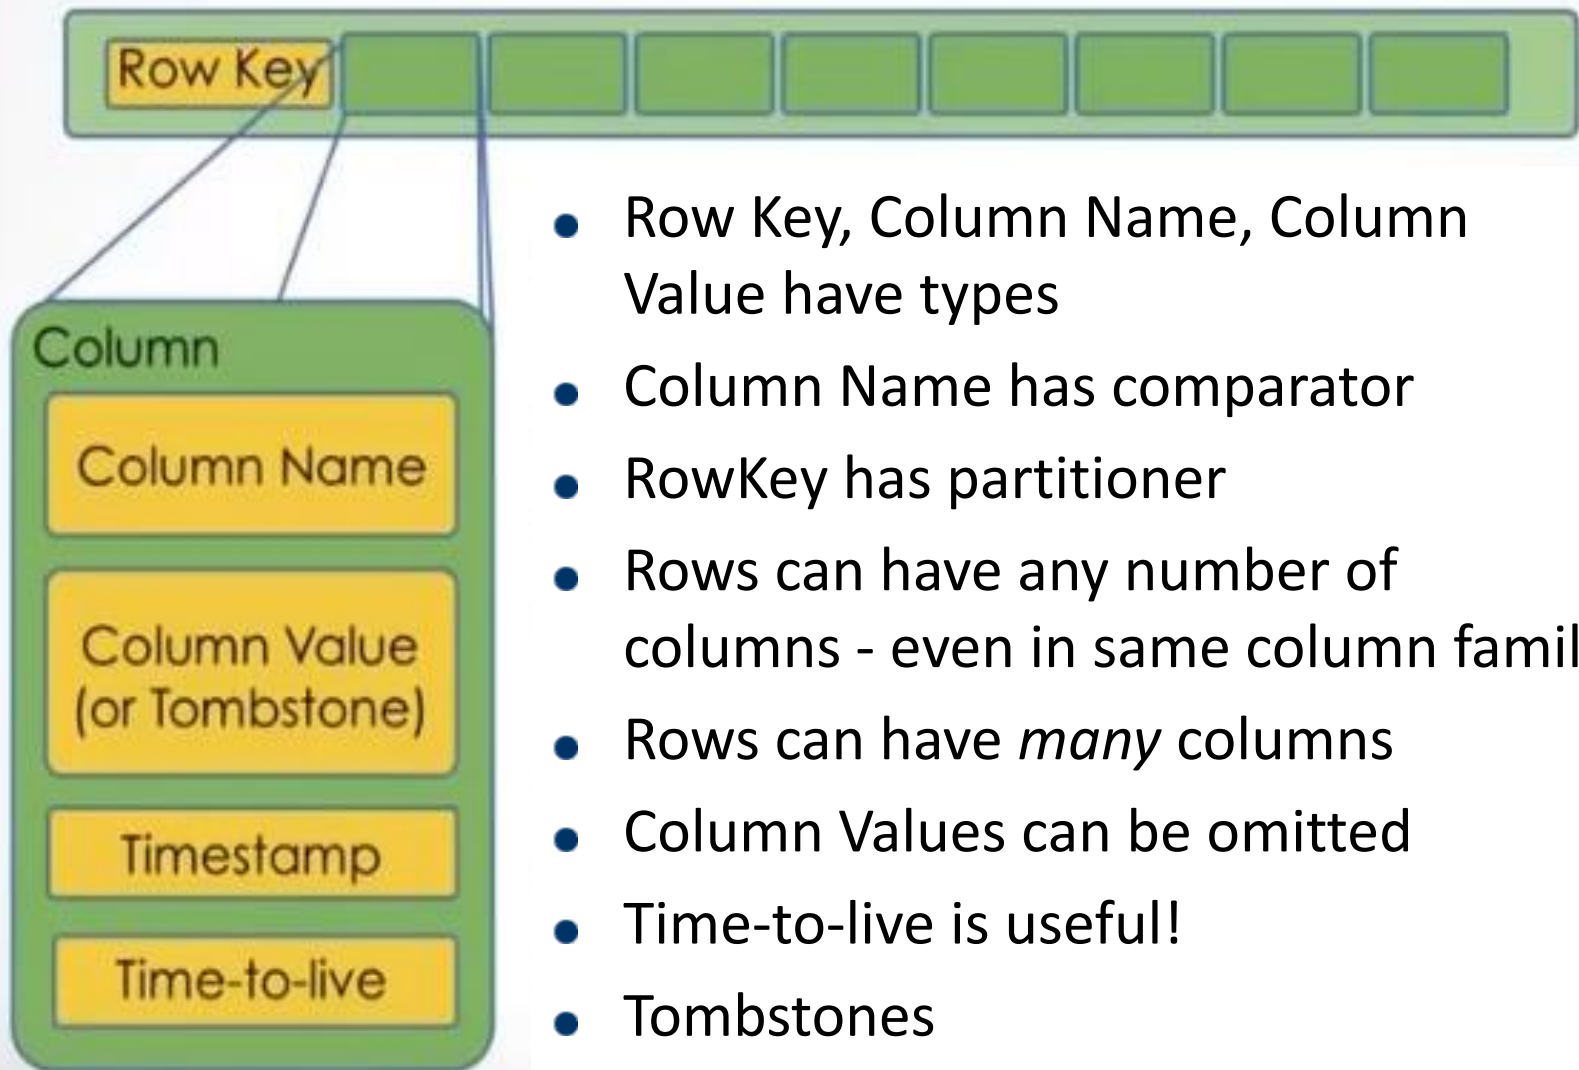
Column Family



Column Family

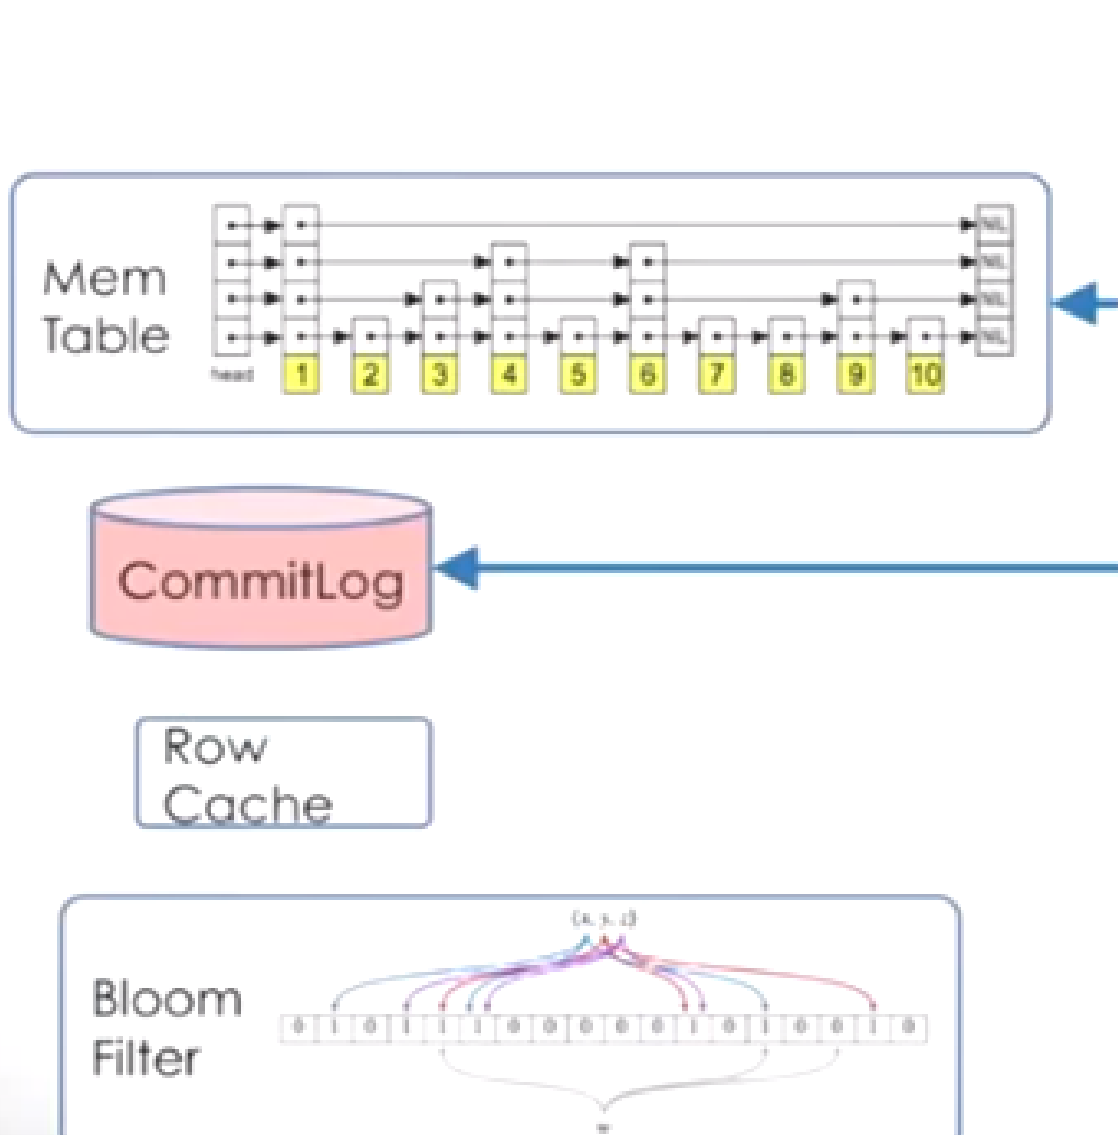


# C\* Data Model (Contd.)





- Row Key, Column Name, Column Value have types
- Column Name has comparator
- RowKey has partitioner
- Rows can have any number of columns - even in same column family
- Rows can have *many* columns
- Column Values can be omitted
- Time-to-live is useful!
- Tombstones

# C\* Data Model: Writes





- Insert into MemTable
- Dump to CommitLog
- No read
- Very Fast!
- Blocks on CPU before O/I!

# C\* Data Model: Reads



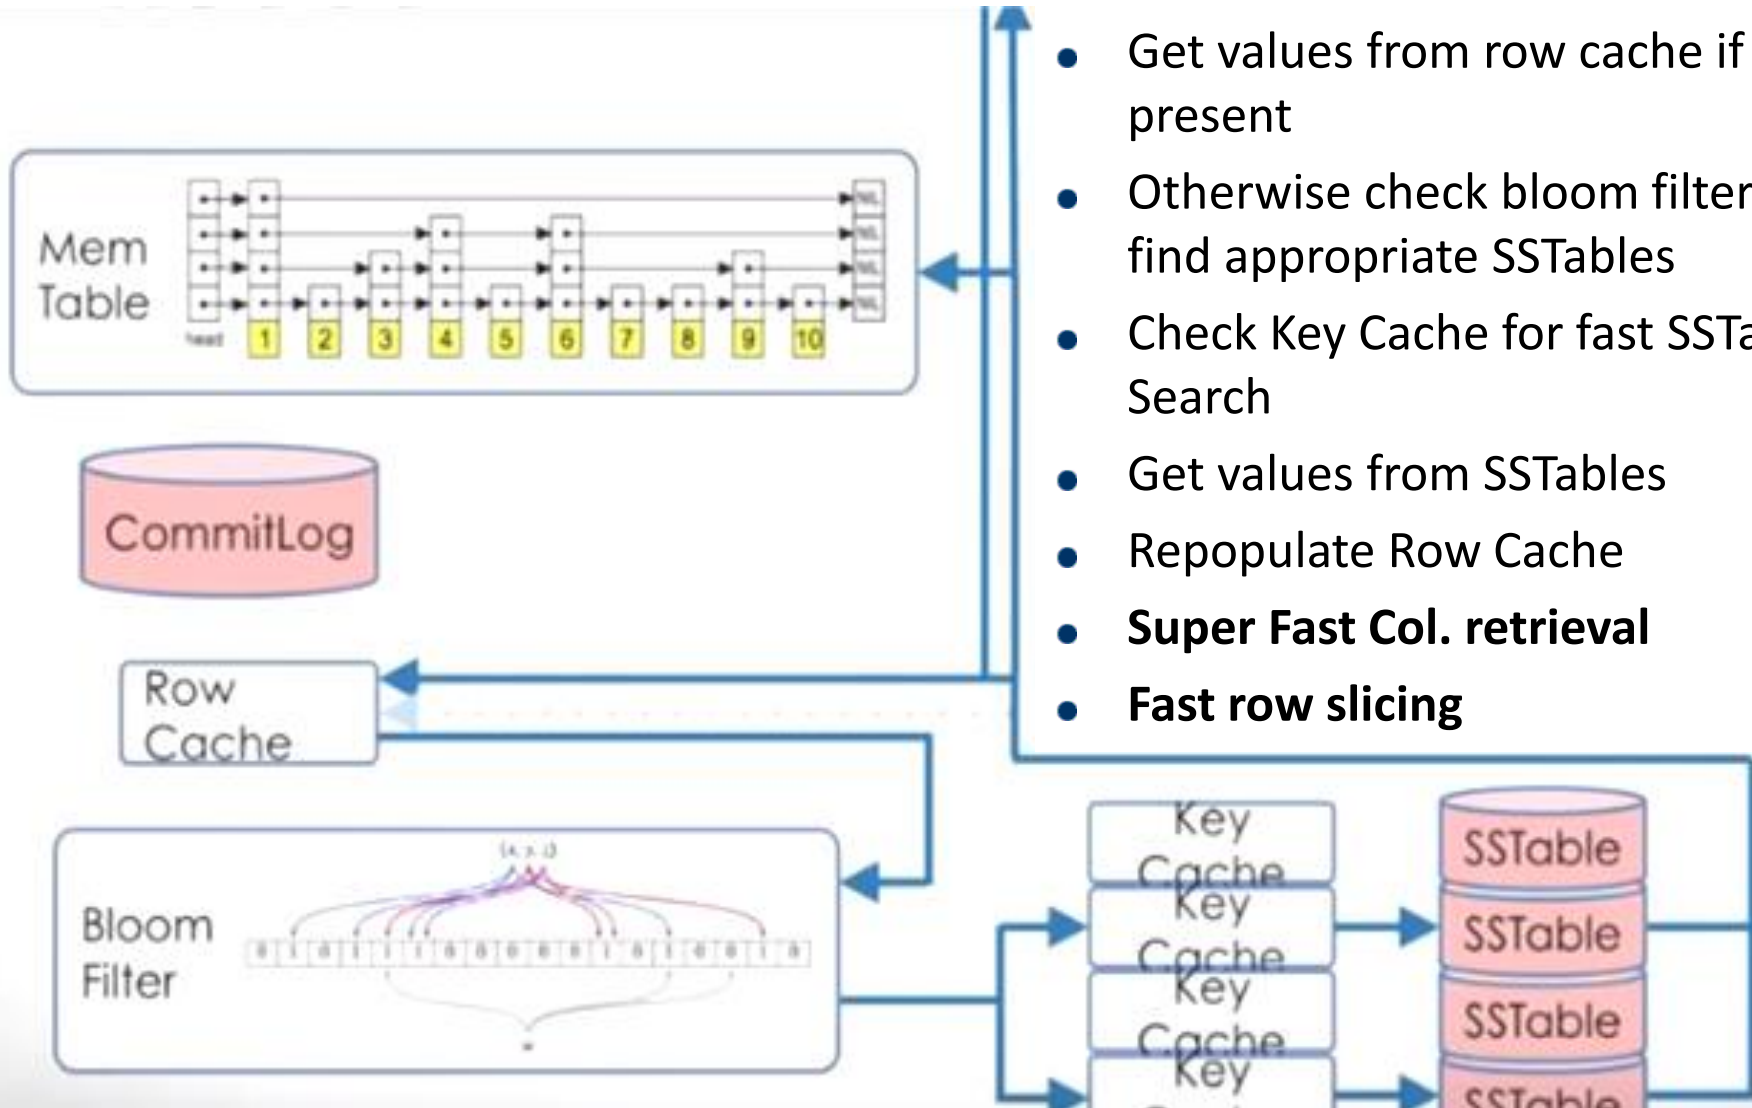
- Get values from Memtable
- Get values from row cache if present
- Otherwise check bloom filter to find appropriate SSTables
- Check Key Cache for fast SSTable Search
- Get values from SSTables
- Repopulate Row Cache
- **Super Fast Col. retrieval**
- **Fast row slicing**





# C\* Data Model: Reads (Contd.)

- Get values from Memtable
- Get values from row cache if present
- Otherwise check bloom filter to find appropriate SSTables
- Check Key Cache for fast SSTable Search
- Get values from SSTables
- Repopulate Row Cache
- **Super Fast Col. retrieval**
- **Fast row slicing**



# Introducing CQL

- **CQL is a reintroduction of schema** so that you don't have to read code to understand the data model.
- **CQL creates a common language** so that details of the data model can be easily communicated.
- **CQL is a best-practices Cassandra interface** and hides the messy details.

# Introducing CQL (Contd.)

```
CREATE TABLE users (  
  id uuid PRIMARY KEY,  
  lastname varchar,  
  firstname varchar,  
  dateOfBirth timestamp );
```

```
INSERT INTO users (id,lastname, firstname, dateofbirth)  
VALUES (now(),'Berryman','John','1975-09-15');
```

```
UPDATE users SET firstname = 'John'  
WHERE id = f74c0b20-0862-11e3-8cf6-b74c10b01fc6;
```

```
SELECT dateofbirth,firstname,lastname FROM users ;
```

dateofbirth	firstname	lastname
1975-09-15 00:00:00-0400	John	Berryman

# Remember this:

- Cassandra finds rows fast
- Cassandra scans columns fast
- Cassandra *does not* scan rows

# The CQL/Cassandra Mapping

```
CREATE TABLE employees (  
  name text PRIMARY KEY,  
  age int,  
  role text  
);
```

name | age | role

-----+-----+-----

john | 37 | dev

eric | 38 | ceo

	age	role
john	37	dev

	age	role
eric	38	ceo

# The CQL/Cassandra Mapping

```
CREATE TABLE employees (
  company text,
  name text,
  age int,
  role text,
  PRIMARY KEY (company,name)
);
```

company | name | age | role

-----+-----+-----+-----

OSC | eric | 38 | ceo

OSC | john | 37 | dev

RKG | anya | 29 | lead

RKG | ben | 27 | dev

RKG | chad | 35 | ops

		eric:age	eric:role	john:age	john:role		
	OSC	38	dev	37	dev		
		anya:age	anya:role	ben:age	ben:role	chad:age	chad:role
RKG		29	lead	27	dev	35	ops

# The CQL/Cassandra Mapping (Contd.)

```
CREATE TABLE example (  
  A text,  
  B text,  
  C text,  
  D text,  
  E text,  
  F text,  
  PRIMARY KEY ((A,B),C,D)  
);
```

A | B | C | D | E | F  
-- + --- + --- + --- + --- + ---

a | b | c | d | e | f

a | b | c | g | h | i

a | b | j | k | l | m

a | n | o | p | q | r

s | t | u | v | w | x

	c:d:E	c:d:F	c:g:E	c:g:F	j:k:E	j:k:F
a:b	e	f	h	i	l	m

	o:p:E	o:p:F
a:n	q	r

	u:v:E	u:v:F
s:t	w	x

# CQL for Sets, Lists and Maps

- Collection Semantics
  - Sets hold list of unique elements
  - Lists hold ordered, possibly repeating elements
  - Maps hold a list of key-value pairs
- Uses same old Cassandra datastructure
- Declaring

```
CREATE TABLE mytable(  
  X text,  
  Y text,  
  myset set<text>,  
  mylist list<int>,  
  mymap map<text, text>,  
  PRIMARY KEY (X,Y)  
);
```

Collection fields  
can not be used  
in primary keys



# Inserting

```
INSERT INTO mytable (row, myset)  
VALUES (123, { 'apple', 'banana'});
```

```
INSERT INTO mytable (row, mylist)  
VALUES (123, ['apple','banana','apple']);
```

```
INSERT INTO mytable (row, mymap)  
VALUES (123, {1:'apple',2:'banana'})
```

# Updating

```
UPDATE mytable SET myset = myset + {'apple','banana'}
```

```
WHERE row = 123;
```

```
UPDATE mytable SET myset = myset - { 'apple' }
```

```
WHERE row = 123;
```

```
UPDATE mytable SET mylist = mylist + ['apple','banana']
```

```
WHERE row = 123;
```

```
UPDATE mytable SET mylist = ['banana'] + mylist
```

```
WHERE row = 123;
```

```
UPDATE mytable SET mymap['fruit'] = 'apple'
```

```
WHERE row = 123
```

```
UPDATE mytable SET mymap = mymap + { 'fruit':'apple' }
```

```
WHERE row = 123
```

# SETS

```
CREATE TABLE mytable(  
  X text,  
  Y text,  
  myset set<int>,  
  PRIMARY KEY (X,Y)  
);
```

```
X | Y | myset  
---+---+-----
```

```
a | b | {1,2}
```

```
a | c | {3,4,5}
```

	b:myset:1	b:myset:2	c:myset:3	c:myset:4	c:myset:5
a					

# LISTS

```
CREATE TABLE mytable(  
  X text,  
  Y text,  
  mylist list<int>,  
  PRIMARY KEY (X,Y)  
);
```

X | Y | mylist

---+---+-----

a | b | [1,2]



	b:mylist:f7e545 <u>00</u> 39..8d	b:mylist:f7e545 <u>01</u> 39..8d
a	1	2

# MAPS

```
CREATE TABLE mytable(  
  X text,  
  Y text,  
  mymap map<text,int>,  
  PRIMARY KEY (X,Y)  
);
```

X | Y | mymap

---+---+-----

a | b | {m:1,n:2}

a | c | {n:3,p:4,q:5}

	b:mymap:m	b:mymap:n	c:mymap:n	c:mymap:p	c:mymap:q
a	1	2	3	4	5

# Example

**(in cqlsh)**

```
CREATE KEYSPACE test WITH replication =  
    {'class': 'SimpleStrategy', 'replication_factor': 1};  
USE test;  
CREATE TABLE stuff ( a int, b int, myset set<int>,  
    mylist list<int>, mymap map<int,int>, PRIMARY KEY (a,b));  
UPDATE stuff SET myset = {1,2}, mylist = [3,4,5], mymap = {6:7,8:9} WHERE a = 0  
AND b = 1;  
SELECT * FROM stuff;
```

**(in cassandra-cli)**

```
use test;  
list stuff ;
```

**(in cqlsh)**

```
SELECT key_aliases,column_aliases from system.schema_columnfamilies WHERE  
keyspace_name = 'test' AND columnfamily_name = 'stuff';
```

# Conclusion

- CQL is a reintroduction of schema
- CQL creates a common data modeling language
- CQL is a best-practices Cassandra interface
- CQL let's you take advantage of the C\* Data structure
- CQL protocol is binary and therefore interoperable with any language
- CQL is asynchronous and fast (Thrift transport layer is synchronous)
- CQL allows the possibility for prepared statements