



Leader election

- Network of processes.
- Want to distinguish exactly one, as the “leader”.
- Eventually, exactly one process should output “leader” (set special status variable to “leader”).
- Motivation: Leader can take charge of:
 - Communication
 - Coordinating data processing (e.g., in commit protocols)
 - Allocating resources
 - Scheduling tasks
 - Coordinating consensus protocols
 - ...




Leader Election: the idea



- We study Leader Election in rings: Depicts basic characteristics of large family of problems in real distributed system (LeLann 1977)
- Illustrates techniques and principles
- Good for lower bounds and impossibility results

The Problem

- Final states of processes partitioned in two classes:




- In every *admissible* execution, exactly one process (the *leader*) enters an elected state. All remaining enter a *non-elected* state.

Ring Networks

- ▶ We assume that the network graph G is a ring consisting of n processes.
- ▶ Numbered $1 \dots n$ in the clockwise direction.
- ▶ We often count $\text{mod } n$, allowing 0 to be another name for process n , $n + 1$ another name for process 1, \dots
- ▶ The processes associated with the nodes of G do not know their indices, nor those of their neighbors.

Processes in a Ring



- ▶ We assume that message-generation and transition functions are defined in terms of local, relative names of their neighbors.



The LCR Algorithm

Algorithm LCR (informal)


Each process sends its identifier around the ring. When a process receives an incoming identifier, it compares that identifier to its own. If the incoming identifier is greater than its own, it keeps passing the identifier; if it is less than its own, it discards the incoming identifier; if it is equal to its own, the process declares itself the leader.

- ▶ Decentralized, Uniform algorithm.
- ▶ Uses only unidirectional communication.
- ▶ Uses only comparison operations on the UIDs.
- ▶ Only the leader performs an output.

Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.


Synchronous Ring



Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.
- ▶ All processes have UIDs
 - ▶ Do now know the UIDs of the other processes.


Synchronous Ring



Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.
- ▶ All processes have UIDs
 - ▶ Do now know the UIDs of the other processes.
- ▶ First round


Synchronous Ring



Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.
- ▶ All processes have UIDs
 - ▶ Do now know the UIDs of the other processes.
- ▶ First round
- ▶ Second round


Synchronous Ring



Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.
- ▶ All processes have UIDs
 - ▶ Do now know the UIDs of the other processes.
- ▶ First round
- ▶ Second round
- ▶ Next rounds


Synchronous Ring



Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.
- ▶ All processes have UIDs
 - ▶ Do now know the UIDs of the other processes.
- ▶ First round
- ▶ Second round
- ▶ Next rounds


Synchronous Ring



Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.
- ▶ All processes have UIDs
 - ▶ Do now know the UIDs of the other processes.
- ▶ First round
- ▶ Second round
- ▶ Next rounds


Synchronous Ring



Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.
- ▶ All processes have UIDs
 - ▶ Do now know the UIDs of the other processes.
- ▶ First round
- ▶ Second round
- ▶ Next rounds


Synchronous Ring



Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.
- ▶ All processes have UIDs
 - ▶ Do now know the UIDs of the other processes.
- ▶ First round
- ▶ Second round
- ▶ Next rounds


Synchronous Ring



Example Execution

- ▶ Let's assume a synchronous ring of $n = 8$ processes.
 - ▶ Processes are indexed from 1 to 8 clockwise.
- ▶ All processes have UIDs
 - ▶ Do now know the UIDs of the other processes.
- ▶ First round
- ▶ Second round
- ▶ Next rounds
- ▶ Leader election – process 2

Synchronous Ring



Variations of the problem

There are several variations of the problem:

- ▶ The ring can be either unidirectional or bidirectional.
- ▶ The number n of nodes may be either known or unknown to the processes.
- ▶ Processes may be identical or can be distinguished by each starting with a **unique identifier** (UID).
- ▶ It might be required that all not-elected processes eventually output the value “non-leader”.
- ▶ It might be required that all non-elected processes eventually output the UID of the leader.
- ▶ We might wish to elect k leaders.
- ▶ ...

Algorithm Properties

Let i_{max} denote the index of the process with the maximum UID, and u_{max} its UID.

- ▶ Process i_{max} is elected leader at the end of round n .
- ▶ No other processes apart from i_{max} ends up in “elected” state.
- ▶ The time complexity is $\mathcal{O}(n)$
- ▶ The message complexity varies...
 - ▶ $\mathcal{O}(n^2)$ — worst case,
 - ▶ $\mathcal{O}(n)$ — best case,
 - ▶ $\mathcal{O}(n \log n)$ — average case.

Algorithm FloodMax

Every process maintains a record of the maximum UID it has seen so far (initially its own). At each round, each process propagates this maximum on all of its outgoing edges. After $diam(G)$ rounds, if the maximum value seen is the process's own UID, the process elects itself the leader; otherwise, it is a non-leader.

- ▶ Processes are not aware of the total number of processes (n).
- ▶ Processes are aware of the network diameter — $\delta = diam(G)$
- ▶ Comparison-based algorithm.


Pseudo-code for FloodMax

```
#DEFINE UID = <...>;
#DEFINE  $\delta$  = <...>;
void main() {
    bool leader = false;
    int max_id = UID;
    for (int i = 0 ; i <  $\delta$ ; i++ ) {
        sendMessage(max_id);
        while (int new_msg = readMessage()) {
            if (new_msg > max_id)
                max_id = new_msg;
        }
    }
    if (max_id == UID)
        leader = true;
}
```

Example of Execution for FloodMax Algorithm

- ▶ Let a synchronous distributed system of $n = 8$ processes..
 - ▶ General network where $\delta = 3$
 - ▶ Processes are index $1 \dots 8$


General Network



Example of Execution for FloodMax Algorithm

- ▶ Let a synchronous distributed system of $n = 8$ processes..
 - ▶ General network where $\delta = 3$
 - ▶ Processes are index $1 \dots 8$
- ▶ The processes have UID.
 - ▶ Not aware of the UID of the other processes.


General Network



Example of Execution for FloodMax Algorithm

- ▶ Let a synchronous distributed system of $n = 8$ processes..
 - ▶ General network where $\delta = 3$
 - ▶ Processes are index $1 \dots 8$
- ▶ The processes have UID.
 - ▶ Not aware of the UID of the other processes.
- ▶ First Round


General Network



Example of Execution for FloodMax Algorithm

- ▶ Let a synchronous distributed system of $n = 8$ processes..
 - ▶ General network where $\delta = 3$
 - ▶ Processes are index $1 \dots 8$
- ▶ The processes have UID.
 - ▶ Not aware of the UID of the other processes.
- ▶ First Round
- ▶ Second Round


General Network



Example of Execution for FloodMax Algorithm

- ▶ Let a synchronous distributed system of $n = 8$ processes..
 - ▶ General network where $\delta = 3$
 - ▶ Processes are index $1 \dots 8$
- ▶ The processes have UID.
 - ▶ Not aware of the UID of the other processes.
- ▶ First Round
- ▶ Second Round
- ▶ Leader Election

General Network



Properties of FloodMax Algorithm

Let n processes and m channels, where the process with the highest UID is i_{max} .

- ▶ Process i_{max} is elected leader at the end of round δ .
- ▶ No other process is in state “elected”.
- ▶ Time complexity is $\mathcal{O}(\text{diam}(G))$.
- ▶ Message complexity $\mathcal{O}(\text{diam}(G) \cdot m)$.