

---

# Consensus


---

In a distributed system,  
agreement among multiple  
processes on a single data  
value, despite failures.

Once they reach a decision  
on a value, that decision is  
final.

# Why Consensus?

# Consensus



Generalizes to  $N$  nodes/processes.

# The Distributed Consensus Problem

We assume  $n$  processes, connected by a synchronous, undirected graph where each process has a unique ID.

Each process  $u$  receives an input value  $i_u$  from the set  $S$ , that is  $i_u \in S$ .

An algorithm solves the problem of distributed consensus if it adheres to the following specifications:

1. **Agreement:** No pair of processes agrees on different output values, that is,  $\nexists u, v : o_u \neq o_v$
2. **Validity:** If all processes start with the same value  $i \in S$ , i.e.,  $\forall u \in [1, n] : i_u = i$ , then value  $i$  is the only possible decision value, that is  $\forall u \in [1, n] : o_u = i$
3. **Termination:** All processes eventually decide.

## SimpleConsensus Algorithm

Each process  $u \in [1, n]$  maintains a list  $l_u$  with pairs of **IDs** and **input values**. Initially the list contains only one set: the ID of  $u$  and the input value  $i_u \in S$ . In each round, all processes transmit the list  $l$  to their local neighborhood. When they receive list  $l_v$  from a neighbor  $v$ , they merge it with their internal list. After  $\delta + 1$  rounds, all processes maintain a list containing a pair  $(u, i_u)$  for each other process of the system. Then they apply a predefined **consensus rule** and terminate by outputting the common value  $o \in S$ .


- ▶ Each process knows  $\delta$ .
- ▶ The algorithm solves the consensus problem.
- ▶ The **consensus rule** can be: minimum value, average value, majority ...

## Example of execution of SimpleConsensus

Let a synchronous network of  $n = 6$  processes and  $\delta = 2$ .

- Consensus rule: simple majority

### General Graph




## Example of execution of SimpleConsensus

Let a synchronous network of  $n = 6$  processes and  $\delta = 2$ .

- Consensus rule: simple majority

1st Round – message transmission






## Example of execution of SimpleConsensus

Let a synchronous network of  $n = 6$  processes and  $\delta = 2$ .

- Consensus rule: simple majority

1st Round – processing






## Example of execution of SimpleConsensus

Let a synchronous network of  $n = 6$  processes and  $\delta = 2$ .

- Consensus rule: simple majority

2nd Round – message transmission




## Example of execution of SimpleConsensus

Let a synchronous network of  $n = 6$  processes and  $\delta = 2$ .

- Consensus rule: simple majority

2nd Round – processing




## Example of execution of SimpleConsensus

Let a synchronous network of  $n = 6$  processes and  $\delta = 2$ .

- Consensus rule: simple majority

3rd Round – message transmission




## Example of execution of SimpleConsensus

Let a synchronous network of  $n = 6$  processes and  $\delta = 2$ .

- Consensus rule: simple majority

3rd Round – processing






## Example of execution of SimpleConsensus

Let a synchronous network of  $n = 6$  processes and  $\delta = 2$ .

- Consensus rule: simple majority

3rd Round– decision



# Properties of SimpleConsensus Algorithm

Let a synchronous network  $G$  with  $n$  processes and  $m$  channels

- ▶ At the end of round  $\delta$  each process  $u \in [1, n]$  will maintain a list  $l_u = \{(1, i_1), (2, i_2), \dots, (n, i_n)\}$
- ▶ The lists maintained by all processes are identical, i.e.,  $\forall u \in [1, n] : l_u = l$
- ▶ The time complexity is  $\mathcal{O}(\text{diam}(G))$
- ▶ The message complexity is  $\mathcal{O}(\text{diam}(G) \cdot m)$
- ▶ The message complexity in bits is  $\mathcal{O}(\text{diam}(G) \cdot n \cdot m)$

## Considerations

How will the execution evolve if failures occur during the transmission of messages ?

Given the presence of failures,

- ▶ can we guarantee the correctness of SimpleConsensus ?
- ▶ can we identify failure ?
- ▶ can we prevent/deal with failure ?



# Fundamental limitation

## Theorem

*Let  $G$  be the graph constituting of nodes 1 and 2 connected by a single edge. Then, there is no algorithm that solves the coordinated attack problem on  $G$  given an unbounded number of link failures.*

- ▶ Impossible to solve basic consensus problems when dealing with totally unreliable network.
- ▶ To overcome, it is necessary to strengthen the model
  - ▶ Assume an upper bound on the number of link failures.
  - ▶ Assume that link failures occur with a probability  $p$ .
- ▶ or relax the problem requirements
  - ▶ Allow the possibility of violating the agreement condition.
  - ▶ Allow the possibility of violating the validity condition.
- ▶ Allow processes to use randomization.



## Stopping Failures


Processes may simply stop arbitrarily without warning, at any point during a round of execution of a distributed algorithm. The process will halt immediately and terminate without further interaction with the other processes of the system.

- ▶ Stopping failures model unpredictable processor crashes.
- ▶ We assume an upper bound  $\sigma$  on the number of stopping failures
  - ▶ such an upper bound holds for the complete execution of the distributed system.
  - ▶ is equivalent to other measures, e.g., rate of stopping failure per round.

## Example of execution of FloodSet algorithm

Let a synchronous complete graph  $n = 4$  and  $\sigma = 2$ .


### Complete Graph



## Example of execution of FloodSet algorithm

Let a synchronous complete graph  $n = 4$  and  $\sigma = 2$ .


1st Round – process 3 fails



## Example of execution of FloodSet algorithm

Let a synchronous complete graph  $n = 4$  and  $\sigma = 2$ .

2nd Round – process 4 fails






## Example of execution of FloodSet algorithm

Let a synchronous complete graph  $n = 4$  and  $\sigma = 2$ .

3rd Round – no failures



## Example of execution of FloodSet algorithm

Let a synchronous complete graph  $n = 4$  and  $\sigma = 2$ .

3rd Round – agreement

