

# Practical Byzantine Fault Tolerance

## Bibliography

M. Castro and B. Liskov. Practical Byzantine fault tolerance and proactive recovery.

*ACM Trans. Comput. Syst.*, 20:398–461, Nov. 2002.

<http://www.disi.unitn.it/~montreso/ds/papers/PbftTocs.pdf>

# Assumptions

- System model
  - Asynchronous distributed system with  $N$  processes
  - Unreliable channels
- Unbreakable cryptography
  - Message  $m$  is signed by its sender  $i$ , and we write  $\langle m \rangle_{\sigma(i)}$ , through:
    - Public/private key pairs
    - Message authentication codes (MAC)
  - A digest  $d(m)$  of message  $m$  is produced through collision-resistant hash functions

## Specification

- State machine replication
  - Replicated service with a state and deterministic operations operating on it
  - Clients issue a request and block waiting for reply
- Safety
  - The system satisfies linearizability, provided that  $N > 3f + 1$
  - Regardless of “faulty clients”...
    - all operations performed by faulty clients are observed in a consistent way by non-faulty clients
  - The algorithm does not rely on synchrony to provide safety...
- Liveness
  - It relies on synchrony to provide liveness
  - Assumes  $delay(t)$  does not grow faster than  $t$  indefinitely
  - Weak assumption – if network faults are eventually repaired
  - Circumvent the impossibility results of FLP

# Assumptions

- Failure model
  - Up to  $f$  Byzantine servers
  - $N > 3f$  total servers
  - (Potentially Byzantine clients)
- Independent failures
  - Different implementations of the service
  - Different operating systems
  - Different root passwords, different administrator

## Theorem

To tolerate up to  $f$  malicious nodes,  $N$  must be equal to  $3f + 1$

## Proof

- It must be possible to proceed after communicating with  $N - f$  replicas, because the faulty replicas may not respond
- But the  $f$  replicas not responding may be just slow, so  $f$  of those that responded might be faulty
- The correct replicas who responded ( $N - 2f$ ) must outnumber the faulty replicas, so

$$N - 2f > f \Rightarrow N > 3f$$


- So,  $N > 3f$  to ensure that at least a correct replica is present in the reply set
- $N = 3f + 1$ ; more is useless
  - more and larger messages
  - without improving resiliency

## Processes and views

- Replicas IDs:  $0 \dots N - 1$
- Replicas move through a sequence of configurations called **views**
- During view  $v$ :
  - **Primary** replica is  $i$ :  $i = v \bmod N$
  - The other are **backups**
- **View changes** are carried out when the primary appears to have failed

## The algorithm

- To invoke an operation, the client sends a request to the primary
- The primary multicasts the request to the backups
- Quorums are employed to guarantee ordering on operations
- When an order has been agreed, replicas execute the request and send a reply to the client
- When the client receives at least  $f + 1$  identical replies, it is satisfied



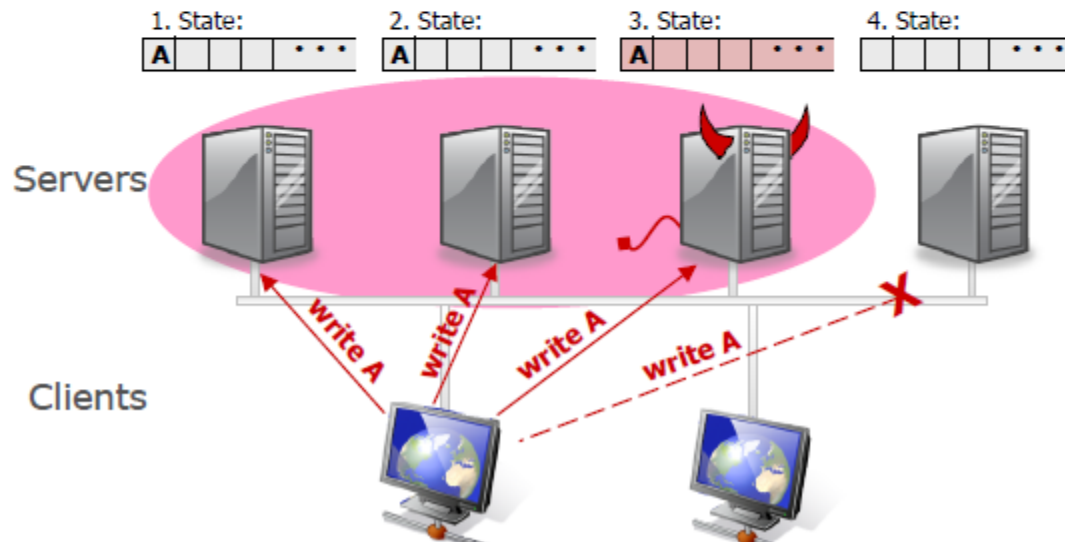
## Problems

- The primary could be faulty!
  - could ignore commands; assign same sequence number to different requests; skip sequence numbers; etc
  - backups monitor primary's behavior and trigger view changes to replace faulty primary
- Backups could be faulty!
  - could incorrectly store commands forwarded by a correct primary
  - use dissemination Byzantine quorum systems
- Faulty replicas could incorrectly respond to the client!
  - Client waits for  $f + 1$  matching replies before accepting response



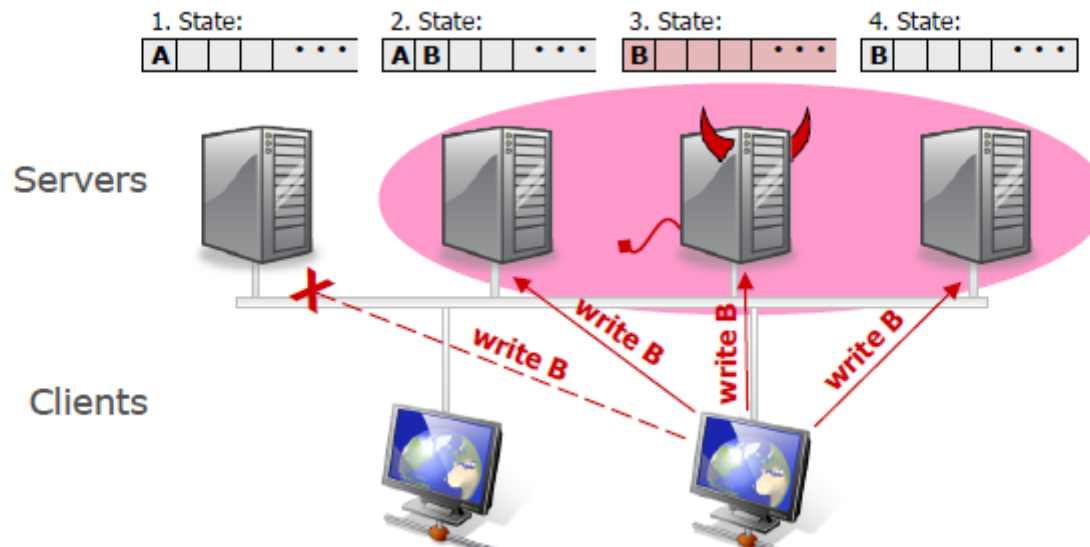
## The general idea

- Algorithm steps are justified by **certificates**
  - Sets (quorums) of signed messages from distinct replicas proving that a property of interest holds
- With quorums of size at least  $2f + 1$ 
  - Any two quorums intersect in at least one correct replica
  - There is always one quorum that contains only non-faulty replicas



## The general idea

- Algorithm steps are justified by **certificates**
  - Sets (quorums) of signed messages from distinct replicas proving that a property of interest holds
- With quorums of size at least  $2f + 1$ 
  - Any two quorums intersect in at least one correct replica
  - There is always one quorum that contains only non-faulty replicas




## Protocol schema

- Normal operation
  - How the protocol works in the absence of failures
  - hopefully, the common case
- View changes
  - How to depose a faulty primary and elect a new one
- Garbage collection
  - How to reclaim the storage used to keep certificates
- Recovery
  - How to make a faulty replica behave correctly again (not here)

## State

- The internal state of each of the replicas include:
  - the state of the actual service
  - a message log containing all the messages the replica has accepted
  - an integer denoting the replica current view


## Client request



$\langle \text{REQUEST}, o, t, c \rangle_{\sigma(c)}$

- $o$ : state machine operation
- $t$ : timestamp (used to ensure exactly-once semantics)
- $c$ : client id
- $\sigma(c)$ : client signature

## Pre-prepare phase



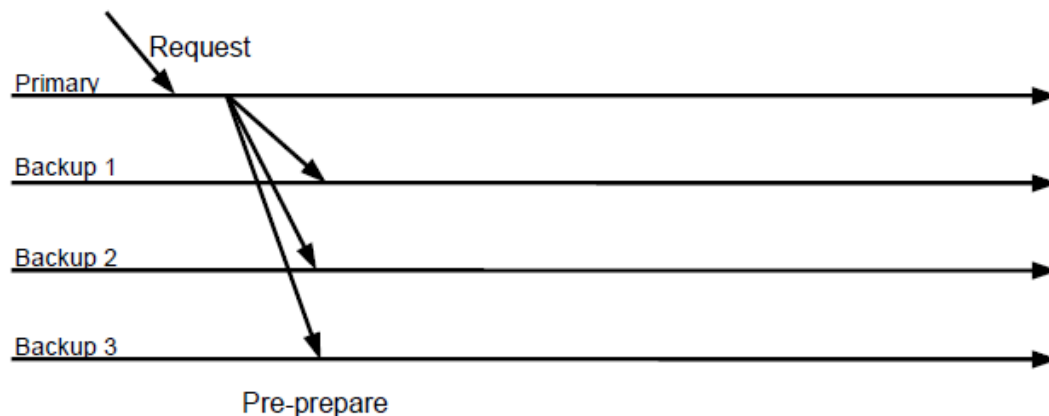
$$\langle \langle \text{PRE-PREPARE}, v, n, d(m) \rangle_{\sigma(p)}, m \rangle$$

- $v$ : current view
- $n$ : sequence number
- $d(m)$ : digest of client message
- $\sigma(p)$ : primary signature
- $m$ : client message


# Pre-prepare phase

$\langle \langle \text{PRE-PREPARE}, v, n, d(m) \rangle_{\sigma(p)}, m \rangle$

- Correct replica  $i$  accepts PRE-PREPARE if:
  - the PRE-PREPARE message is well-formed
  - the current view of  $i$  is  $v$
  - $i$  has not accepted another PRE-PREPARE for  $v, n$  with a different digest
  - $n$  is between two water-marks  $L$  and  $H$   
(to avoid sequence number exhaustion caused by faulty primaries)
- Each accepted PRE-PREPARE message is stored in the accepting replica's message log (including the primary's)
- Non-accepted PRE-PREPARE messages are just discarded



## Prepare phase



$\langle \text{PREPARE}, v, n, d(m) \rangle_{\sigma(i)}$

- Accepted by correct replica  $j$  if:
  - the PREPARE message is well-formed
  - current view of  $j$  is  $v$
  - $n$  is between two water-marks  $L$  and  $H$
- Replicas that send PREPARE accept the sequence number  $n$  for  $m$  in view  $v$
- Each accepted PREPARE message is stored in the accepting replica's message log




## Prepare certificate (P-certificate)

- Replica  $i$  produces a **prepare certificate**  $\text{prepared}(m, v, n, i)$  iff its log holds:
  - The request  $m$
  - A PRE-PREPARE for  $m$  in view  $v$  with sequence number  $n$
  - Log contains  $2f$  PREPARE messages from different backups that match the PRE-PREPARE
- $\text{prepared}(m, v, n, i)$  means that a quorum of  $(2f + 1)$  replicas agrees with assigning sequence number  $n$  to  $m$  in view  $v$

### Theorem

There are no two non-faulty replicas  $i, j$  such that  $\text{prepared}(m, v, n, i)$  and  $\text{prepared}(m', v, n, j)$ , with  $m \neq m'$

## Commit phase




$\langle \text{COMMIT}, v, n, d(m), i \rangle_{\sigma(i)}$

- After having collected a P-certificate  $\text{prepared}(m, v, n, i)$ , replica  $i$  sends a COMMIT message
- Accepted if:
  - The COMMIT message is well-formed
  - Current view of  $i$  is  $v$
  - $n$  is between two water-marks  $L$  and  $H$

## Commit certificate (C-Certificate)

- **Commit certificates** ensure total order across views
  - we guarantee that we can't miss prepare certificates during a view change
- A replica has a certificate **committed**( $m, v, n, i$ ) if:
  - it had a P-certificate **prepared**( $m, v, n, i$ )
  - log contains  $2f + 1$  matching COMMIT from different replicas (possibly including its own)
- Replica executes a request after it gets commit certificate for it, and has cleared all requests with smaller sequence numbers

## Reply phase



$\langle \text{REPLY}, v, t, c, i, r \rangle_{\sigma(i)}$

- $r$  is the reply
- Client waits for  $f + 1$  replies with the same  $t, r$
- If the client does not receive replies soon enough, it broadcast the request to all replicas

## View change

- A un-satisfied replica backup  $i$  **mutinies**:
  - stops accepting messages (except VIEW-CHANGE and NEW-VIEW)
  - multicasts  $\langle \text{VIEW-CHANGE}, v + 1, P, i \rangle_{\sigma(i)}$
  - $P$  contains a P-certificate  $P_m$  for each request  $m$   
(up to a given number, see garbage collection)
- Mutiny succeeds if the new primary collects a **new-view certificate**  $V$ :
  - a set containing  $2f + 1$  VIEW-CHANGE messages
  - indicating that  $2f + 1$  distinct replicas (including itself) support the change of leadership

## View change

The “**primary elect**”  $p'$  (replica  $v + 1 \bmod N$ ):

- extracts from the new-view certificate  $V$  the highest sequence number  $h$  of any message for which  $V$  contains a P-certificate
- creates a new PRE-PREPARE message for any client message  $m$  with sequence number  $n \leq h$  and add it to the set  $O$ 
  - if there is a P-certificate for  $n, m$  in  $V$

$$O \leftarrow O \cup \langle \text{PRE-PREPARE}, v + 1, n, d_m \rangle_{\sigma(p')}$$

- Otherwise

$$O \leftarrow O \cup \langle \text{PRE-PREPARE}, v + 1, n, d_{\text{null}} \rangle_{\sigma(p')}$$

- $p'$  multicasts  $\langle \text{NEW-VIEW}, v + 1, V, O \rangle_{\sigma(p')}$

## View change

- Backup accepts a  $\langle \text{NEW-VIEW}, v + 1, V, O \rangle_{\sigma(p')}$  message for  $v + 1$  if
  - it is signed properly by  $p'$
  - $V$  contains valid VIEW-CHANGE messages for  $v + 1$
  - the correctness of  $O$  can be locally verified (repeating the primary's computation)
- Actions:
  - Adds all entries in  $O$  to its log (so did  $p'$ !)
  - Multicasts a PREPARE for each message in  $O$
  - Adds all PREPARES to the log and enters new view

## Garbage collection

- A correct replica keeps in log messages about request  $o$  until:
  - $o$  has been executed by a majority of correct replicas, and
  - this fact can be proven during a view change
- Truncate log with stable checkpoints
  - Each replica  $i$  periodically (after processing  $k$  requests) checkpoints state and multicasts  $\langle \text{CHECKPOINT}, n, d, i \rangle$ 
    - $n$ : last executed request
    - $d$ : state digest
- A set  $S$  containing  $2f + 1$  equivalent CHECKPOINT messages from distinct processes are a proof of the checkpoint's correctness  
(stable checkpoint certificate)



## View Change, revisited

- Message  $\langle \text{VIEW-CHANGE}, v + 1, n, S, C, P, i \rangle_{\sigma(i)}$ 
  - $n$ : the sequence number of the last stable checkpoint
  - $S$ : the last stable checkpoint
  - $C$ : the checkpoint certificate ( $2f + 1$  checkpoint messages)
- Message  $\langle \text{NEW-VIEW}, v + 1, n, V, O \rangle_{\sigma(p')}$ 
  - $n$ : the sequence number of the last stable checkpoint
  - $V, O$ : contains only requests with sequence number larger than  $n$

## Optimizations

- Reducing replies
  - One replica designated to send reply to client
  - Other replicas send digest of the reply
- Lower latency for writes (4 messages)
  - Replicas respond at Prepare phase (tentative execution)
  - Client waits for  $2f + 1$  matching responses
- Fast reads (one round trip)
  - Client sends to all; they respond immediately
  - Client waits for  $2f + 1$  matching responses

## Optimizations: cryptography


- Reducing overhead
  - Public-key cryptography only for view changes
  - MACs (message authentication codes) for all other messages
- To give an idea (Pentium 200Mhz)
  - Generating 1024-bit RSA signature of a MD5 digest: 43ms
  - Generating a MAC of the same message:  $10\mu s$

# Talk Overview

- Problem
- Assumptions
- Algorithm
- Implementation
- Performance
- Conclusions

# Algorithm Properties


- Arbitrary replicated service
  - complex operations
  - mutable shared state
- Properties (safety and liveness):
  - system behaves as correct centralized service
  - clients eventually receive replies to requests
- Assumptions:
  - $3f+1$  replicas to tolerate  $f$  Byzantine faults (optimal)
  - strong cryptography
  - **only for liveness:** eventual time bounds



# Algorithm Overview

## State machine replication:

- deterministic replicas start in same state
- replicas execute same requests in same order
- correct replicas produce identical replies




**Hard: ensure requests execute in same order**

# Ordering Requests

## Primary-Backup:


- View designates the primary replica



- Primary picks ordering
- Backups ensure primary behaves correctly
  - certify correct ordering
  - trigger view changes to replace faulty primary

# Quorums and Certificates

quorums have at least  $2f+1$  replicas



quorums intersect in at least one correct replica

- **Certificate**  $\equiv$  set with messages from a quorum
- Algorithm steps are justified by certificates



# Algorithm Components

- Normal case operation
- View changes
- Garbage collection
- Recovery


All have to be designed to work together

# Normal Case Operation

- Three phase algorithm:
  - *pre-prepare* picks order of requests
  - *prepare* ensures order within views
  - *commit* ensures order across views
- Replicas remember messages in log
- Messages are authenticated
  - $\langle \bullet \rangle_{\sigma_k}$  denotes a message sent by k

# Pre-prepare Phase


assign sequence number  $n$  to request  $m$  in view  $v$



backups accept pre-prepare if:

- in view  $v$
- never accepted pre-prepare for  $v, n$  with different request


# Prepare Phase



# Order Within View


**No *P-certificates* with the same view and sequence number and different requests**

If it were false:



**one correct replica in common  $\Rightarrow m = m'$**

# Commit Phase



Request  $m$  executed after:

- having *C-certificate*( $m, v, n$ )
- executing requests with sequence number less than  $n$


# View Changes

- Provide liveness when primary fails:
  - timeouts trigger view changes
  - select new primary ( $\equiv$  view number mod  $3f+1$ )
- But also need to:
  - preserve safety
  - ensure replicas are in the same view long enough
  - prevent denial-of-service attacks

# View Change Safety

**Goal: No *C-certificates* with the same sequence number and different requests**


- Intuition: if replica has *C-certificate*( $m, v, n$ ) then





# View Change Protocol

**send *P*-certificates:**  $\langle \text{VIEW-CHANGE}, v+1, \mathbf{P}, 2 \rangle_{\sigma_2}$



***X*-certificate:**  $\langle \text{NEW-VIEW}, v+1, \mathbf{X}, \mathbf{O} \rangle_{\sigma_1}$

**pre-prepares matching**  
***P*-certificates with highest views in  $\mathbf{X}$**


- pre-prepare for  $m, v+1, n$  in new-view
- Backups multicast prepare messages for  $m, v+1, n$

backups multicast prepare messages for pre-prepares in  $\mathbf{O}$

# Garbage Collection

Truncate log with **certificate**:

- periodically checkpoint state (**K**)
- multicast  $\langle \text{CHECKPOINT}, h, D(\text{checkpoint}), i \rangle_{\sigma_i}$
- all collect  $2f+1$  checkpoint messages  
 **$S\text{-certificate}(h, \text{checkpoint})$**



discard messages and checkpoints



reject messages

send  **$S\text{-certificate}$**  and checkpoint in view-changes